

Automated Risk of Bias Assessment in Systematic Reviews Using Large Language Models: A Multi-Model Single-Agent Benchmark on 227 Randomised Controlled Trials

Luiz Felipe Guain Teixeira^{1,2,3} Cassiano Pereira de Barros⁴
Renato Vicente^{1,2,3}

¹TELUS Digital Research Hub

²Center for Artificial Intelligence and Machine Learning, University of São Paulo

³Institute of Mathematics, Statistics and Computer Science, University of São Paulo

⁴Paulista School of Medicine, Federal University of São Paulo

{luiz.felipe.teixeira, rvicente}@usp.br

Abstract

Background: Risk of bias (RoB) assessment is a mandatory but labour-intensive step in systematic review production. Large language models (LLMs) offer the potential to automate this task in a zero-shot setting, yet systematic benchmarks spanning multiple frontier models and inference strategies are lacking.

Methods: Eight state-of-the-art LLMs (Gemini 2.5 Pro, Gemini 2.5 Flash, GPT-5.4, Claude Haiku 4.5, DeepSeek v4-Flash, Grok 4.3, Qwen 3.7-Plus, and GLM-5)

were evaluated on zero-shot, single-agent RoB 1.0 assessment of 227 randomised controlled trials drawn from 127 Cochrane systematic reviews published in 2024. Models were accessed via the OpenRouter API in June 2026 using a fixed prompt and provider-default generation parameters (exact OpenRouter identifier strings in Table 2). Ground-truth annotations were taken from the published Cochrane consensus judgements across seven methodological domains (1,589 cells total). Performance was assessed using accuracy, Macro-F1, and Gwet's AC1, with study-clustered bootstrap 95% confidence intervals. Two models (Gemini 2.5 Flash and DeepSeek v4-Flash) were additionally evaluated under a single-call inference mode for comparison with the per-domain strategy.

Results: Overall accuracy ranged from 57.6% to 62.5% across models, a spread smaller than the width of individual 95% confidence intervals; Macro-F1 confidence intervals overlapped across all eight models. AC1 provided the clearest separation: Claude Haiku 4.5 ($AC1 = 0.485$; 95% CI $[0.443, 0.526]$) had non-overlapping CIs relative to Gemini 2.5 Pro ($[0.344, 0.420]$) and GPT-5.4 ($[0.349, 0.426]$), though this comparison is unadjusted and exploratory. Unclear Risk was the most consistently under-detected class (recall < 0.20 in three of seven domains across most models). Consensus failures (cells where no model was correct) accounted for only 14.2% of the benchmark, confirming substantial non-redundancy among models; the remaining 85.8% of cells had at least one correct model. The comparison of inference modes yielded small, largely non-significant differences: ten of twelve Δ metrics had confidence intervals spanning zero, and D7 (Other Bias)—the domain hypothesised to benefit most from joint assessment—showed negligible change under either mode.

Conclusions: Frontier LLMs can perform zero-shot RoB 1.0 assessment at accuracy levels broadly consistent with estimates of human inter-rater variability for this task [1], though direct comparison on this specific corpus was not performed; no single model dominates across all metrics, and no consistent advantage was found for single-call over per-domain inference.

46 The primary practical distinction between inference modes is computational:
47 single-call reduces inference calls by $7\times$ without a reliable accuracy trade-off.
48 All analyses are exploratory; results should be interpreted as directional
49 benchmarks rather than definitive rankings.
50 **Keywords:** risk of bias; systematic review; large language models; natural
51 language processing; Cochrane; evidence synthesis; automation

1 Introduction

Systematic reviews and meta-analyses occupy the apex of the evidence hierarchy in evidence-based medicine [2]. Their validity depends critically on the quality of the primary studies included: a meta-analysis that aggregates trials with uncontrolled confounding or selective outcome reporting inherits—and can amplify—those methodological flaws. The Cochrane Risk of Bias 1.0 tool (RoB 1.0) [3] was developed precisely to render this quality appraisal explicit and reproducible. By requiring reviewers to judge each included trial across seven pre-specified methodological domains, RoB 1.0 transforms a subjective quality impression into a structured, auditable record.

Despite this formalisation, RoB assessment remains a major bottleneck in systematic review production. Each trial must be read by two independent reviewers, who compare judgements and resolve disagreements through discussion or adjudication. For a review covering dozens of trials, this process routinely demands tens to hundreds of person-hours. Inter-rater reliability, even among trained reviewers, is moderate at best [4, 5]—a consequence of genuine ambiguity in reporting, domain-specific expertise requirements, and the epistemic character of the *Unclear Risk* category, which denotes insufficient information rather than a graded risk level. These limitations have motivated sustained interest in automating, or at least augmenting, the RoB assessment workflow.

Early automation efforts relied on rule-based systems and supervised machine learning trained on sentence-level text fragments [6]. Such approaches demonstrated that natural language processing could identify methodological descriptions in trial reports with above-chance accuracy, but they were constrained by limited contextual understanding, sensitivity to reporting style, and the need for substantial labelled training data for each domain.

Large language models (LLMs) trained on diverse corpora with instruction-following capabilities represent a qualitatively different paradigm [7, 8]. Without task-specific fine-tuning, general-purpose LLMs can be prompted to read

81 a full trial report, locate domain-relevant passages, and produce structured
82 judgements with free-text justifications—closely mimicking the expert reviewer
83 workflow. A prompt template can address all seven RoB 1.0 domains either
84 sequentially—one domain per inference call—or simultaneously in a single call;
85 whether this choice affects assessment quality is an open empirical question.
86 As frontier LLMs have grown more capable and their inference costs have fallen,
87 evaluating their effectiveness for RoB assessment has become both feasible
88 and practically important. A recent preregistered comparative study applied
89 four contemporary LLMs to 100 RCTs from Cochrane reviews under the RoB 1.0
90 instrument, reporting interobserver κ of 0.27–0.39 and concluding that none
91 reached the reliability threshold required for autonomous use [9]. Whether this
92 shortfall is uniform across the RoB 1.0 structure—or whether reliable assistance
93 is already viable for specific domains and intervention types while systematic fail-
94 ures persist elsewhere—has direct implications for how LLMs can be integrated
95 into real screening workflows.

96 The present study addresses this question with a systematic comparison
97 of eight state-of-the-art LLMs on zero-shot, single-agent RoB 1.0 assessment
98 applied to 227 randomised controlled trials drawn from 127 Cochrane systematic
99 reviews published in 2024. All models were evaluated on a common benchmark
100 under identical conditions: the same prompt, the same fully text-based study
101 representations produced by a standardised PDF-to-Markdown pipeline, and the
102 same ground-truth human annotations. Beyond aggregate accuracy and Macro-F1,
103 we report Cohen’s κ , Gwet’s AC1, and per-class F1—placing κ and AC1 side by side
104 to expose reliability estimates that diverge under class imbalance. Domain- and
105 intervention-stratified analyses identify where systematic biases arise and locate
106 the improvement opportunities most likely to yield practical gains for assisted
107 screening. Additionally, we compare two inference strategies—per-domain and
108 single-call—to assess whether querying all seven domains simultaneously affects
109 performance relative to sequential per-domain calls.

110 2 Methods

111 2.1 Dataset

112 The ground-truth dataset was constructed from Cochrane systematic reviews
113 published in 2024. Of the 149 reviews published that year, 127 were accessible
114 through our institutional subscription; from each accessible review, up to three
115 included randomised controlled trials (RCTs) were selected on the basis of full-
116 text PDF availability: all included RCTs for which a PDF was retrievable were
117 included, up to a maximum of three per review, yielding a final corpus of **227**
118 **RCTs** drawn from 127 distinct Cochrane reviews.

119 Risk of bias annotations were obtained directly from the published Cochrane
120 review reports and reflect the consensus judgements produced by the original
121 review teams following standard Cochrane procedures (independent assessment
122 by at least two trained reviewers, with disagreements resolved by discussion or
123 arbitration). These annotations were not re-produced or independently adjudi-
124 cated for the present study; they represent the published Cochrane consensus,
125 which constitutes the reference standard used by the field. Although the revised
126 Risk of Bias 2.0 tool (RoB 2; Sterne et al. 10) is now recommended for new system-
127 atic reviews, it reorganises the domain structure and changes the rating scale,
128 rendering it incompatible with the existing RoB 1.0 corpus. RoB 1.0 remains the
129 basis for the large majority of existing systematic review annotations — including
130 all 127 reviews sampled here and most of the large labelled datasets used for
131 training and evaluation of automated systems — so this study’s focus on RoB 1.0
132 is both corpus-driven and practically motivated. Each study was annotated using
133 the Cochrane Risk of Bias 1.0 (RoB 1.0) tool [3], which covers seven methodolog-
134 ical domains: D1 (random sequence generation), D2 (allocation concealment),
135 D3 (blinding of participants and personnel), D4 (blinding of outcome assessment),
136 D5 (incomplete outcome data), D6 (selective outcome reporting), and D7 (other
137 sources of bias). Each domain received one of three judgements—*Low Risk*, *High*

138 *Risk, or Unclear Risk*—yielding 1,589 annotated cells (227×7).

139 Class prevalence varied substantially across domains (Table 1). Low Risk was
140 the largest class in six of seven domains, and the majority class ($>50\%$) in four
141 of them (D1, D5, D6, D7); only D3 (blinding of participants and personnel) had
142 High Risk as its largest class (45%), reflecting the predominance of unblinded
143 pharmacological trials in the corpus. At the skewed extreme, D1 had a 69%
144 Low Risk majority and only 5 High Risk cases (2%); D5 and D7 were similarly
145 imbalanced at 63% and 61% Low Risk, respectively. D4 (blinding of outcome
146 assessors) was the most balanced domain overall (44% Low Risk, 28% High Risk,
147 29% Unclear Risk), closely followed by D2, in which Low Risk and Unclear Risk
148 were nearly equally represented (50% vs. 46%).

Table 1: Ground-truth class distribution per RoB domain ($n = 227$ studies per domain). LR = Low Risk; HR = High Risk; UR = Unclear Risk.

Domain	LR	HR	UR	Majority class	Majority (%)
D1	157	5	65	Low Risk	69.2
D2	113	10	104	Low Risk	49.8
D3	74	103	50	High Risk	45.4
D4	99	63	65	Low Risk	43.6
D5	143	39	45	Low Risk	63.0
D6	125	25	77	Low Risk	55.1
D7	139	31	57	Low Risk	61.2
Total	850 (53.5%)	276 (17.4%)	463 (29.1%)	Low Risk	53.5

149 **Intervention type.** Each study was assigned a single intervention category:
150 drug/medication ($n = 104$), non-pharmacological intervention ($n = 75$), proce-
151 dure ($n = 30$), or medical device ($n = 18$).

152 **Clinical domain.** Studies were tagged with one or more of 24 Cochrane Re-
153 view Group categories (multi-label; 120 of 227 studies belonged to more than
154 one group). The most frequent categories were Child Health ($n = 72$), Heart &
155 Circulation ($n = 45$), and Kidney Disease ($n = 31$); the least frequent were Skin
156 Disorders ($n = 3$), Effective Practice & Health Systems ($n = 4$), Rheumatology
157 ($n = 5$), and Wounds ($n = 5$). For stratified analyses, each multi-domain study

158 was included in every applicable group.




Figure 1: Distribution of the 227 RCTs by intervention type.




Figure 2: Distribution of the 227 RCTs by clinical domain (Cochrane Review Group categories, multi-label; total exceeds 227).

159 2.2 Document Preparation Pipeline

160 Full-text study reports (PDF) were processed through a two-stage pipeline that
161 converts each article into a standardised plain-text representation before it
162 is submitted to any model. Converting to a common format ensures that all

163 eight models receive identical textual input, so that observed differences in
164 performance reflect model behaviour rather than differences in how each model
165 internally parses or prioritises content from a raw PDF.

166 **Stage 1 — PDF-to-Markdown conversion.** Each PDF was converted to struc-
167 tured Markdown using MinerU [11], which preserves headings, tables, and inline
168 text whilst replacing embedded figures with placeholder tokens of the form
169 [Figure k].

170 **Stage 2 — Figure description.** Each figure placeholder was resolved by Gem-
171 ini 2.5 Flash (google/gemini-2.5-flash, accessed via OpenRouter in May–June 2026),
172 a vision-language model (VLM) that received the corresponding page image
173 together with surrounding article context and generated a structured textual
174 description following a fixed eight-field schema (figure type, axes and units, leg-
175 end and groups, visual summary, plotted patterns, and unreadable or uncertain
176 areas; full prompt reproduced in Supporting Information S10). The description
177 was then inserted in place of the placeholder. After Stage 2, the document was
178 a self-contained Markdown file in which all information—including any figures
179 present in the source document—was represented as plain text. The result is
180 a self-contained Markdown file in which the full article content—text, tables,
181 and figures—is represented uniformly as plain text, forming the common input
182 submitted to all eight models.

183 **2.3 Risk of Bias Assessment**

184 All models were evaluated in a zero-shot, single-agent setting: each processed
185 Markdown document was submitted to the model with a fixed assessment prompt
186 (reproduced in full in Supporting Information S1), which specified the domain
187 criteria, the three admissible judgements (Low Risk, High Risk, Unclear Risk), and
188 the required JSON output format. No few-shot examples were provided, and the
189 prompt was used without modification across all models.

190 **Inference modes.** We evaluated two inference strategies that differ in how

the seven RoB domains are requested from the model. In the *per-domain* mode, seven sequential inference calls are made per study: each call supplies the full Markdown text together with a single domain identifier (e.g. Domain: D1), and the model returns a JSON object containing a judgement and justification for that domain only. In the *single-call* mode, a single inference call per study is made: the model receives the full Markdown text and is instructed to return a JSON array containing judgements and justifications for all seven domains simultaneously. All eight models were evaluated under the per-domain mode. Per-domain mode was adopted as the primary protocol because each RoB 1.0 domain represents an independent methodological criterion; separating calls ensures that each domain is assessed under a dedicated instruction context, without the model being required to simultaneously manage seven distinct evaluation tasks in a single generation pass. Gemini 2.5 Flash and DeepSeek v4-Flash were additionally evaluated under the single-call mode on the same 227 studies, enabling a controlled comparison of the two inference strategies whilst holding model, prompt content, and corpus constant. These two models were selected because their published list prices per token on OpenRouter were the lowest among the eight models evaluated, making a second full-corpus run economically practical. The results of this comparison are reported in Section 3.5.

2.4 Language Models

Eight frontier LLMs were evaluated (Table 2), accessed via the OpenRouter API in June 2026 using the identifier strings listed therein. No model-specific fine-tuning, prompt engineering beyond the fixed assessment prompt (Section 2.3), or retrieval augmentation was applied.

Evaluation protocol. All calls used provider-default generation parameters and each study-condition pair received a single inference draw with no output aggregation, constituting a hyperparameter-free single-draw evaluation protocol. This design reflects two constraints. First, decoding parameters such as

temperature interact with model architecture and training objective in model-family-specific ways; varying them across models would confound generation configuration with prompt design, rendering cross-model comparisons uninterpretable. Second, post-hoc reliability strategies—including self-consistency decoding [12], output ensembling, and majority voting over repeated draws—are optimisations orthogonal to native zero-shot capability and constitute a separate experimental question outside the scope of this benchmark.

Reproducibility and versioning. OpenRouter routes each identifier string to the provider-current checkpoint at inference time; underlying weights may be updated by the provider without a corresponding change to the public identifier. All results are therefore conditioned on the checkpoint states active during the June 2026 inference window and should not be assumed to generalise to subsequent revisions of the same identifier strings.

Table 2: Language models evaluated. All models accessed via the OpenRouter API in June 2026 using the identifier strings shown; zero-shot single-agent mode throughout; all 227 studies evaluated per model. Context length: maximum tokens accepted per call as reported by OpenRouter at access time. Release date: date the model became available on OpenRouter. Exact checkpoint versions are opaque at the provider level; the identifiers shown are the strings passed to the API and represent the model state at the time of inference.

Model	OpenRouter identifier	Context	Release
Gemini 2.5 Pro	google/gemini-2.5-pro	1M	Jun 2025
Gemini 2.5 Flash	google/gemini-2.5-flash	1M	Jun 2025
GPT-5.4	openai/gpt-5.4	1M	Mar 2026
Claude Haiku 4.5	anthropic/claude-haiku-4.5	200K	Oct 2025
DeepSeek v4-Flash	deepseek/deepseek-v4-flash	1M	Apr 2026
Grok 4.3	x-ai/grok-4.3	1M	Apr 2026
Qwen 3.7-Plus	qwen/qwen3.7-plus	1M	May 2026
GLM-5	z-ai/glm-5	203K	Feb 2026

2.5 Evaluation Metrics

Accuracy. The proportion of study-domain cells in which the model’s judgement matched the ground-truth label, equivalent to the observed agreement p_o used

in the chance-corrected coefficients below. Let $\mathbf{N} = (N_{ij})$ denote the $K \times K$ confusion matrix (here, $K = 3$), with rows indexed by predicted class i and columns indexed by ground-truth class j :

$$\text{Accuracy} = p_o = \frac{\sum_{k=1}^K N_{kk}}{N} = \frac{\sum_{k=1}^K N_{kk}}{\sum_{i=1}^K \sum_{j=1}^K N_{ij}}, \quad (1)$$

where N_{ij} is the number of cells with ground-truth class j predicted as class i , and $N = \sum_{i=1}^K \sum_{j=1}^K N_{ij}$ is the total number of evaluated cells ($N = 1,589$ overall; $N = 227$ per domain; smaller for intervention-type and clinical-domain strata).

Per-class F1. RoB assessment is a three-class problem ($K = 3$). Per-class metrics are computed under a one-vs-rest scheme: each study-domain cell is treated as a binary classification (“class k ” vs. “not class k ”) for each class $k \in \{\text{Low Risk, High Risk, Unclear Risk}\}$, yielding four counts:

- $\text{TP}_k = N_{kk}$: model and ground truth both assigned class k ;
- $\text{FN}_k = \sum_{i \neq k} N_{ik}$: ground truth was k , model assigned another class (positive missed);
- $\text{FP}_k = \sum_{j \neq k} N_{kj}$: model assigned k , ground truth was another class (false alarm);
- $\text{TN}_k = \sum_{i \neq k} \sum_{j \neq k} N_{ij}$: neither model nor ground truth assigned class k .

Precision (positive predictive value) and recall (sensitivity) for class k are

$$\text{Precision}_k = \frac{\text{TP}_k}{\text{TP}_k + \text{FP}_k}, \quad \text{Recall}_k = \frac{\text{TP}_k}{\text{TP}_k + \text{FN}_k}, \quad (2)$$

and their harmonic mean is the F1 score for class k :

$$\text{F1}_k = \frac{2 \cdot \text{Precision}_k \cdot \text{Recall}_k}{\text{Precision}_k + \text{Recall}_k}. \quad (3)$$

Per-class F1 reveals class-specific failure modes independently of overall performance.

255 **Macro-F1.** Per-class F1 is computed over the full pool of 1,589 cells; the three
 256 class-level F1 scores are then averaged with equal weight:

$$\text{Macro-F1} = \frac{1}{K} \sum_{k=1}^K \text{F1}_k, \quad K = 3. \quad (4)$$

257 Unlike Accuracy, Macro-F1 assigns equal weight to each category regardless
 258 of prevalence, penalising models that achieve high accuracy by concentrating
 259 predictions on the majority class.

260 **Cohen's κ .** A chance-corrected agreement coefficient that adjusts observed
 261 agreement p_o for the agreement expected by chance under the assumption that
 262 annotators label independently according to their observed marginal distribu-
 263 tions:

$$\kappa = \frac{p_o - p_e^\kappa}{1 - p_e^\kappa}, \quad p_e^\kappa = \sum_{k=1}^K \hat{p}_{k,\text{model}} \cdot \hat{p}_{k,\text{gt}}, \quad (5)$$

264 where $\hat{p}_{k,\text{model}} = \frac{1}{N} \sum_{j=1}^K N_{kj}$ and $\hat{p}_{k,\text{gt}} = \frac{1}{N} \sum_{i=1}^K N_{ik}$ are the marginal proportions
 265 of class k in the model predictions and ground truth, respectively (row and column
 266 sums of N divided by N). κ is reported alongside AC1 to expose their divergence
 267 under class imbalance, illustrating the prevalence paradox (see Gwet's AC1 below).

268 **Gwet's AC1.** A chance-corrected agreement coefficient robust to the preva-
 269 lence paradox [13, 14]: when marginal distributions are highly skewed, classical
 270 chance-corrected coefficients approach zero even when observed agreement is
 271 substantial. AC1 is defined as

$$\text{AC1} = \frac{p_o - p_e}{1 - p_e}, \quad p_e = \frac{1}{K-1} \sum_{k=1}^K \pi_k (1 - \pi_k), \quad (6)$$

272 where p_o is the observed proportion of agreement, $K = 3$ is the number of
 273 categories, and $\pi_k = (\hat{p}_{k,\text{model}} + \hat{p}_{k,\text{gt}})/2$ is the mean marginal proportion for
 274 category k , with $\hat{p}_{k,\text{model}}$ and $\hat{p}_{k,\text{gt}}$ the marginal proportions of class k in the model
 275 predictions and ground truth, respectively. All metrics—including AC1—are com-
 276 puted globally by pooling all 1,589 study-domain cells; domain-level values

are not aggregated. Given the pronounced class imbalance documented in Table 1—six of seven domains have Low Risk as the majority class—AC1 is more appropriate than Cohen’s κ as the primary agreement metric for this task, since κ approaches zero under heavy prevalence skew even when observed agreement is substantial (the prevalence paradox; Feinstein and Cicchetti 14). Importantly, the ground-truth annotations used here reflect published Cochrane consensus judgements rather than a second independent annotator; AC1 therefore quantifies model-to-published-standard agreement, not inter-rater reliability in the classical sense.

Error direction (LR-anchor framework). Because the three RoB classes are not ordinally exchangeable—Low Risk acts as the “safe” default that passes a study through without additional scrutiny, while High Risk and Unclear Risk both trigger closer examination—we classify non-diagonal errors relative to Low Risk. A *lenient* error occurs when the model assigns Low Risk and the ground truth is High Risk or Unclear Risk (false pass). A *cautious* error occurs when the model assigns High Risk or Unclear Risk and the ground truth is Low Risk (false flag). A *lateral* error occurs when the model confuses High Risk with Unclear Risk or vice versa. Because both non-Low-Risk categories already trigger reviewer scrutiny, lateral errors do not affect the binary pass/no-pass decision; their consequence is second-order, affecting the precision of the label (definite bias vs. insufficient information) rather than whether the study receives attention. This LR-anchor decomposition distinguishes the two first-order error directions: lenient errors risk including potentially biased studies without scrutiny (inflating evidence-base credibility), while cautious errors risk discrediting studies with genuinely low risk of bias (reducing the available evidence base). Both carry clinical costs in opposite directions; their relative severity depends on review context and the scarcity of eligible studies.

304 **2.6 Consensus Failures and Model Complementarity**

305 For each study–domain cell, we recorded whether at least one of the eight
306 models matched the ground truth. Cells where *no* model was correct are termed
307 *consensus failures* and represent cases in which all evaluated models failed
308 simultaneously, regardless of which model was consulted. The *any-correct rate*
309 (proportion of cells where at least one model succeeded) characterises the
310 degree of complementarity among models: a high any-correct rate alongside low
311 individual-model accuracy indicates that different models fail on different cells,
312 yielding non-redundant coverage across the ensemble.

313 **2.7 Stratified Analyses**

314 All primary metrics were recomputed within subgroups defined by intervention
315 type (four categories) and clinical domain (24 Cochrane Review Groups, multi-
316 label). For clinical domain, studies were included in every group to which they
317 belonged, so group sample sizes sum to more than 227.

318 **2.8 Statistical Inference**

319 Analyses are primarily descriptive and exploratory; two post-hoc formal tests
320 were additionally conducted to characterise pairwise performance differences
321 and systematic class-distribution bias. Confidence intervals are reported to char-
322 acterise estimation uncertainty and facilitate comparison across models and
323 conditions; they are not interpreted as the outcome of pre-specified confirma-
324 tory hypothesis tests. All 95% CIs were computed by study-clustered bootstrap
325 (5,000 iterations): the 227 studies were resampled with replacement, all seven
326 domain cells of each sampled study were retained, and the metric of interest
327 was recomputed on the resulting pseudo-sample; the 2.5th and 97.5th percentiles
328 of the bootstrap distribution define the interval. Resampling at the study level
329 preserves within-study correlation across domains (see Supporting Informa-

tion S7 for a comparison with Wilson score intervals). No familywise error rate correction was applied: the number of comparisons across models, domains, classes, and inference modes is large and the outcomes are correlated, making conventional multiplicity corrections either too conservative (Bonferroni) or difficult to calibrate without a pre-specified primary endpoint. Where 95% bootstrap CIs exclude zero, this is noted descriptively (“CI excludes zero”); the 95% level serves as a conventional descriptive reference threshold and carries no confirmatory implication in this exploratory context—it marks an unadjusted signal that may warrant further investigation, not a principled significance claim.

McNemar test. Pairwise differences in overall accuracy between models were assessed using the two-sided continuity-corrected McNemar test [15], applied to the 2×2 table of discordant cells for each model pair. Because all eight models evaluated the same 1,589 cells, errors are paired by design; the McNemar test exploits this within-cell correlation, yielding greater statistical power than comparisons based on marginal confidence intervals.

Stuart–Maxwell test. Systematic deviations in the marginal class distribution of each model’s predictions relative to the ground truth were assessed using the Stuart–Maxwell test of marginal homogeneity [16, 17], which extends McNemar’s test to $k \times k$ square contingency tables ($k = 3$: Low Risk, High Risk, Unclear Risk). The test statistic is $Q = \mathbf{d}^\top \mathbf{S}^{-1} \mathbf{d} \sim \chi^2(k-1)$, where $\mathbf{d}$ is the vector of marginal differences and $\mathbf{S}$ is its covariance matrix. Eight tests were conducted (one per model); p-values were adjusted using the Benjamini–Hochberg false discovery rate procedure [18] at $\alpha = 0.05$.

3 Results

3.1 Overall Performance

Figure 3 summarises accuracy, Macro-F1, per-class F1, and Gwet’s AC1 for all eight models across the 1,589 annotated cells, together with study-clustered

bootstrap 95% confidence intervals (5,000 iterations, resampling the 227 studies—the unit of clustering—with replacement to preserve within-study correlation across the seven domains). Overall accuracy ranged from 57.6% (Gemini 2.5 Pro) to 62.5% (Claude Haiku 4.5), a spread of less than five percentage points. Macro-F1 confidence intervals similarly overlap across all eight models, precluding reliable rank ordering based on this metric alone. By contrast, Gemini 2.5 Flash and DeepSeek v4-Flash tied for the highest Macro-F1 (0.558) while Claude Haiku 4.5—the accuracy leader—ranked sixth (0.541). The per-class breakdown reveals that the Macro-F1 gap is driven primarily by High Risk and Unclear Risk detection: Claude Haiku 4.5 achieved the highest Low Risk F1 (0.743) but the lowest High Risk F1 (0.469) and Unclear Risk F1 (0.413) among all models. Among the three metrics, Gwet’s AC1 shows the clearest separation at the extremes: the 95% CI for Claude Haiku 4.5 ($AC1 = 0.485$; CI [0.443, 0.526]) does not overlap with those of Gemini 2.5 Pro ([0.344, 0.420]) or GPT-5.4 ([0.349, 0.426]), suggesting that Haiku’s higher chance-corrected agreement is unlikely to be attributable to sampling variation alone.

Metric divergence. The discrepancy between accuracy and Macro-F1 rankings reflects differential handling of the class imbalance documented in Table 1. Figure 3 makes the trade-off explicit: models with higher Low Risk F1 tend to achieve higher accuracy but lower High Risk and Unclear Risk F1, since Low Risk is the majority class in six of seven domains. Macro-F1 penalises this behaviour by weighting each class equally, shifting the ranking toward models with more balanced per-class detection. Inspection of Figure 3 also reveals that the rank orderings for accuracy, AC1, and F1-LR are nearly identical across the eight models. This co-linearity suggests that all three metrics are largely co-determined by Low Risk performance: a model that reliably assigns Low Risk to majority-class cells captures most of the variance in both accuracy and chance-corrected agreement, since Low Risk prevalence (53.5%) dominates both the observed agreement count and the marginal distributions that enter the AC1 correction. Macro-F1

Overall model performance — 95% bootstrap CIs (study-clustered, 5 000 iterations)




Figure 3: Overall performance on RoB 1.0 assessment for the eight models ($n = 227$ studies; 1,589 cells). Each row is one model; columns show accuracy (%), Cohen's κ , Gwet's AC1, and F1 for each risk class (F1-LR, F1-HR, F1-UR). κ and AC1 are placed side by side to expose their divergence under class imbalance (prevalence paradox). Bars represent point estimates; error bars are 95% study-clustered bootstrap confidence intervals (5,000 iterations, resampling 227 studies with replacement). Dashed vertical lines show analytical baselines: *Always LR* (green), *Always HR* (red), *Always UR* (orange), and *Random uniform* (blue dotted); baselines with F1 = 0 or $\kappa = 0$ for a given metric are not drawn. For κ and AC1, the dashed grey line marks zero (chance-level agreement). Macro-F1 is reported in Table SS1. Colours identify models consistently throughout the paper. Full numerical values are reported in Table SS1.

disrupts this co-linearity by treating each class as equally important regardless of prevalence. However, equal class weighting does not reflect the empirical frequency of judgements—High Risk accounts for only 17.4% of cells (Table 1)—so macro-F1 amplifies the contribution of the rarest class beyond its share of the data. A prevalence-weighted F1 would instead converge toward accuracy, whereas macro-F1 should be interpreted as a measure of balanced class coverage: how well a model performs across *all three* nominal categories, independently of how often each arises.

Ground-truth uncertainty and inter-annotator agreement. The benchmark treats published Cochrane annotations as the reference standard, but those annotations carry their own measurement error. Even among formally trained

reviewers applying the same protocol, Hartling et al. [1] reported $\kappa = 0.79$ for D1 but only $\kappa = 0.24$ – 0.37 for the remaining six domains, with agreement between independent consensus pairs falling further to $\kappa = 0.05$ – 0.09 in several domains. Agreement also varies substantially between distinct Cochrane review teams evaluating the same trials [4]. These findings contextualise model performance against a noisy reference standard rather than a noise-free human ceiling, and motivate domain-level rather than aggregate interpretation of results. Figure 4 places model performance in this context by displaying the full 9×9 label-agreement matrix across Ground Truth and all eight models.




Figure 4: Label agreement matrix for Ground Truth and all eight models ($n = 1,589$ cells). Cell (i, j) = percentage of cells in which annotator i and annotator j assigned the *same* label, irrespective of correctness. The diagonal is 100% by definition. The white lines isolate the Ground Truth row and column. Colour scale: 55–100% (Blues).

The Ground Truth row (first row) reproduces the known model accuracies

(57.6%–62.5%), representing agreement between each model and the Cochrane reference. All inter-model cells (the lower-right 8×8 block) exceed these values, ranging from 63.1% (Claude Haiku 4.5 and Gemini 2.5 Pro) to 82.0% (Qwen 3.7-Plus and GLM-5). These pairwise statistics are descriptive: no bootstrap confidence intervals were computed for them. Because cells from the same study share textual context and are correlated across domains, the effective sampling unit is the study ($n = 227$), not the cell ($n = 1,589$); each pairwise proportion therefore carries an approximate 95% uncertainty of ± 6 pp ($SE \approx \sqrt{p(1-p)/227}$, evaluated at $p \approx 0.70$). Differences smaller than this margin should not be over-interpreted; the patterns noted below reflect the overall structure of the matrix rather than fine-grained rankings between adjacent cells. This systematic gap—models agree with each other more than they agree with the ground truth—indicates that the eight models share systematic prediction tendencies that diverge from the Cochrane standard in coordinated ways. It does not imply that the models are collectively more reliable than the ground truth; rather, it reflects shared training priors that cause convergent errors.

The magnitude of the gap is interpretable in light of the known reference noise: if Cochrane reviewer pairs themselves agree at ≈ 60 –79% (the $\kappa = 0.24$ –0.37 range, converted to expected agreement under the observed class distribution), then the inter-model agreement range of 63–82% is broadly comparable to or modestly above human inter-rater reliability on this task. The ground-truth column therefore represents one reference annotator among many possible ones, and the model accuracy figures (57–63%) should be interpreted as lower bounds on the models’ latent capability rather than definitive measurements of their performance against a noise-free standard.

Pairwise model comparisons. Marginal bootstrap confidence intervals for accuracy overlap across all eight models; the greater power of the paired McNemar test—which leverages within-cell correlation across identical items—reveals finer distinctions. Gemini 2.5 Pro and GPT-5.4 form a distinct lower-performing

cluster, each significantly outperformed by four of the six remaining models—
 Gemini 2.5 Flash, Claude Haiku 4.5, Qwen 3.7-Plus, and GLM-5 (all $p \leq 0.014$,
 Benjamini–Hochberg corrected, $m = 28$ pairs)—while remaining statistically
 indistinguishable from each other ($p = 0.956$). Differences involving DeepSeek v4-
 Flash and Grok 4.3 were nominally significant ($p = 0.026$ – 0.040) but did not survive
 BH correction. The full pairwise accuracy differences and corresponding p -values
 are shown in Figures 5 and 6.

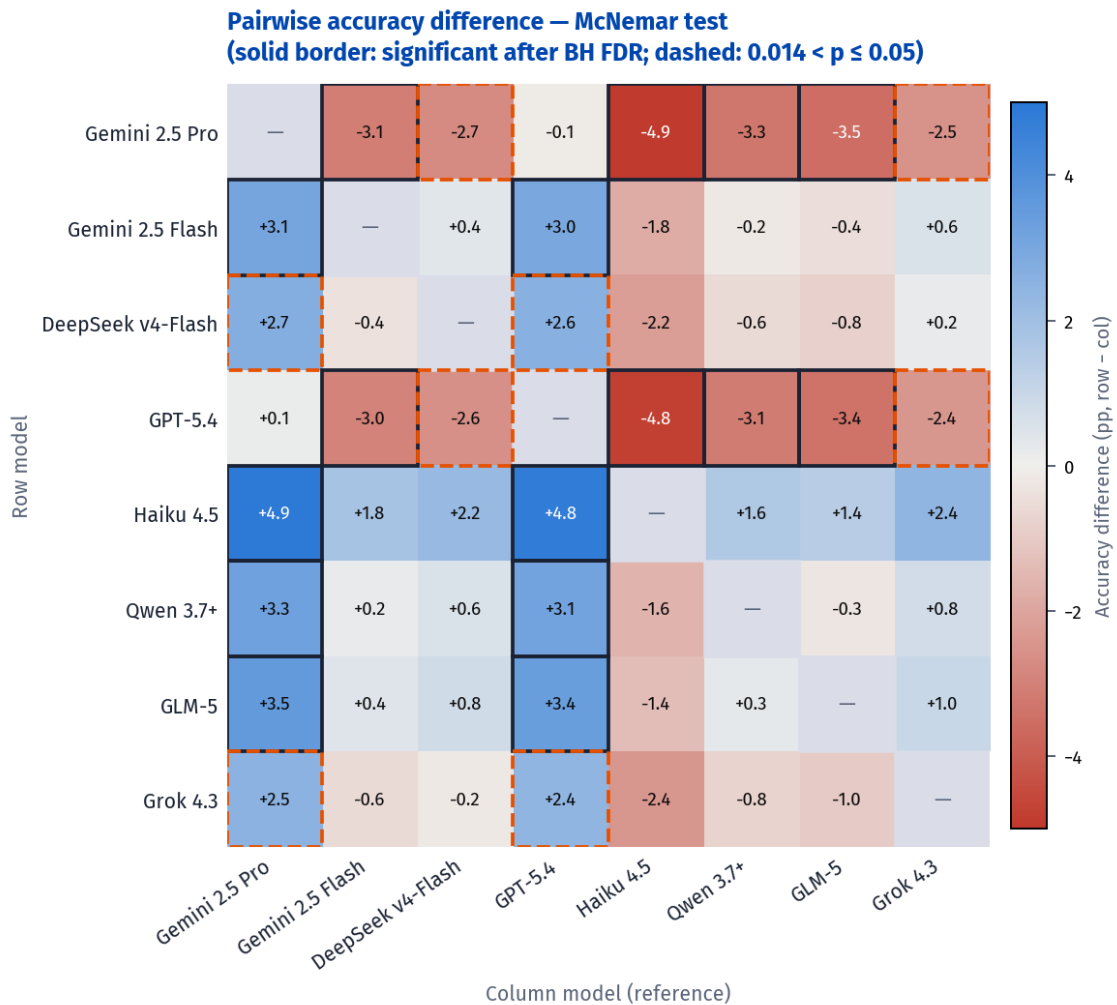


Figure 5: Pairwise accuracy difference matrix (McNemar test). Cell (i, j) = accuracy of the row model minus accuracy of the column model (percentage points), computed over the 1,589 study–domain cells common to all models. Colour scale: diverging blue (row model better) to red (column model better), midpoint at zero. Solid border: pair is significant after Benjamini–Hochberg FDR correction ($\alpha = 0.05$, $m = 28$ pairs); dashed orange border: nominally significant ($p \leq 0.05$) but does not survive BH correction.

Consensus failures and model complementarity. Table 3 summarises the




Figure 6: Pairwise McNemar p -values. Cell (i, j) shows the $-\log_{10}(p)$ for the continuity-corrected McNemar test comparing the row and column models, with the raw p -value annotated. Colour scale: sequential blue (light = $p \approx 1$, dark = $p \ll 0.001$). Solid border: BH-significant; dashed orange: nominally significant ($p \leq 0.05$) but does not survive BH correction.

444 full distribution of per-cell model agreement across the 1,589 study-domain
 445 cells. The distribution is markedly bimodal: the two largest individual categories
 446 are *all models correct* (499 cells, 31.4%) and *consensus failure* (226 cells, 14.2%),
 447 with the intermediate categories together accounting for the remaining 54.4%.
 448 This bimodality suggests that cell difficulty is largely an intrinsic property of the
 449 item rather than the model: most cells are either solvable by all frontier models
 450 evaluated here, or intractable for all of them given the prompt context.

Table 3: Distribution of per-cell model agreement across the 1,589 study-domain cells (8 models). *Majority* and *minority* thresholds are defined relative to $n = 8$; see Section 2.6 for the generalised formula.

Category	N	%
All correct (8/8)	499	31.4
Majority (5–7/8)	418	26.3
Split (4/8)	96	6.0
Minority (1–3/8)	350	22.0
None correct (0/8)	226	14.2
Total	1,589	100.0

451 Stratifying by ground-truth class (Table 4) reveals a clear asymmetry. *Low risk*
452 cells are classified correctly by all eight models in 44.4% of cases, and only 4.9%
453 are consensus failures. *Unclear risk* cells are the hardest: 27.0% are consensus
454 failures—no model assigns the correct label—exceeding even *high risk* (21.4%).
455 This result is consistent with the inherent difficulty of the *unclear risk* category,
456 which requires inferring the *absence* of adequate reporting rather than identifying
457 a positive signal of bias.

Table 4: Per-cell agreement distribution by ground-truth class (row percentages).

Class	n	All (8/8) %	Majority (5–7) %	Split (4) %	Minority (1–3) %	None (0) %
Low risk	850	44.4	25.9	5.1	19.8	4.9
High risk	276	13.8	27.5	9.8	27.5	21.4
Unclear risk	463	18.1	26.3	5.6	22.9	27.0

458 Figure 7 shows the pairwise intersection matrix: cell (i, j) gives $P(A_i \cap A_j)$, the
459 percentage of the 1,589 cells in which *both* model i and model j produced the
460 correct judgement. The diagonal entries report each model’s overall accuracy
461 ($P(A_i)$).




Figure 7: Pairwise intersection matrix. Cell $(i, j) = P(A_i \cap A_j)$: percentage of the 1,589 study-domain cells in which both models were correct. Diagonal entries (bold) = individual model accuracy. Colour scale: 40–65% (YlOrRd); deeper red = greater overlap in correct cells. Upper triangle omitted by symmetry.

Off-diagonal intersection values range from 45.1% (Claude Haiku 4.5 and GPT-5.4) to 53.5% (Qwen 3.7-Plus and GLM-5), all substantially below the respective diagonal entries (57.6%–62.5%). The gap between diagonal and off-diagonal reflects that no two models share all of their correct cells: a non-trivial fraction of each model’s successes are unique to that model. Expressing overlap as the Jaccard similarity $J(A, B) = P(A \cap B) / P(A \cup B)$ —the fraction of cells correct for *at least one* model that are correct for *both*—yields values between 60.0% and 78.1% across the 28 pairs. The highest Jaccard, Qwen 3.7-Plus and GLM-5 ($J = 78.1\%$, $P(A \cap B) = 53.5\%$), indicates that nearly four in five cells that either model answers correctly are answered correctly by both, confirming this as the most redundant pair. The lowest Jaccard, Claude Haiku 4.5 and GPT-5.4 ($J = 60.0\%$, $P(A \cap B) = 45.1\%$), indicates that two in five jointly-correct cells belong exclusively to one model—the highest non-redundancy in the ensemble

475 and a direct predictor of the high complementarity entry visible in Figure 8.

476 Figure 8 shows the full pairwise complementarity matrix: cell (i, j) gives the
 477 percentage of the 1,589 cells in which model i erred and model j was correct.
 478 Values range from 7.4% to 17.4%, with all off-diagonal entries exceeding 7%,
 479 confirming that no two models in the ensemble are redundant.




Figure 8: Pairwise complementarity matrix. Cell (i, j) = percentage of the 1,589 study-domain cells in which the row model erred and the column model was correct. Colour scale: 0–20% (PuOr); purple = high complementarity, orange = low complementarity; grey diagonal = undefined. Full numeric values are reported in Table SS6.

480 The most salient feature of Figure 8 is the Claude Haiku 4.5 column: it is the
 481 darkest column overall, meaning that when any other model errs, Haiku recovers
 482 those cells at the highest rate (12.3%–17.4%). This is the direct complement of
 483 its accuracy profile—Haiku leads on overall accuracy and AC1 by concentrating
 484 predictions on the majority class, which makes its correct cells less likely to be
 485 errors for other models, but also makes its errors structurally distinct from theirs.
 486 The two highest individual entries in the matrix—GPT-5.4 → Haiku (17.4%) and
 487 Gemini 2.5 Pro → Haiku (17.3%)—correspond to the two models with the lowest

488 overall accuracy; their larger error sets overlap least with Haiku's. The same
489 structural distinctiveness is visible in the agreement matrix (Figure 4): Haiku
490 has the lowest mean inter-model agreement across the ensemble (67.9%, versus
491 71.7%–74.9% for all other models), driven by consistently low agreement with
492 Gemini 2.5 Pro and GPT-5.4 (63–64%)—the same dyads that top the complemen-
493 tarity column. High complementarity and low overall agreement are two sides
494 of the same coin: because Haiku assigns labels differently from the ensemble
495 cluster, it agrees with other models less often irrespective of correctness, but
496 those disagreements are disproportionately cases where Haiku is right and the
497 other model is wrong. At the opposite extreme, the Qwen 3.7-Plus ↔ GLM-5 pair
498 shows the lowest bidirectional complementarity (7.4% and 7.6%), indicating the
499 most overlapping error profiles of any pair in the ensemble. Grok 4.3 also shows
500 low complementarity as a *column* model for the Qwen and GLM-5 rows (7.6% and
501 8.9%), suggesting that these three models share a similar error structure that is
502 not well covered by the others. This pattern is reinforced by the agreement matrix
503 (Figure 4): the three highest pairwise inter-model agreements—Qwen 3.7-Plus /
504 GLM-5 (82.0%), GLM-5 / Grok 4.3 (80.2%), and Qwen 3.7-Plus / Grok 4.3 (77.6%)—are
505 all concentrated within this trio; the trio's mean agreement (79.9%) exceeds the
506 mean of cross-trio pairs by 7.2 pp. Taken together, the matrix indicates that
507 Claude Haiku 4.5, Gemini 2.5 Pro, and GPT-5.4 form the most complementary triad:
508 their pairwise entries are uniformly above 10%, and combining any two of them
509 recovers more than a fifth of all cells that either individually misses.

510 **Error direction by model.** When errors are aggregated at the model level,
511 leniency dominates the profiles of Claude Haiku 4.5 and Qwen 3.7-Plus, while
512 Gemini 2.5 Pro and GPT-5.4 show a more mixed pattern consistent with their
513 global high-risk and unclear-risk biases identified by the Stuart–Maxwell test
514 (below). Figure 9 shows the proportional breakdown of lenient, lateral, and
515 cautious errors for each model.

516 The Stuart–Maxwell test examined whether each model's overall predicted




Figure 9: Overall error direction by model under the LR-anchor framework. Each bar shows the proportion of errors classified as lenient (red), lateral (orange), or cautious (blue); bar thickness scales linearly with the total number of errors for that model ($n = 596\text{--}674$). Models are ordered top to bottom as in Figure 3.

class distribution differed significantly from the ground-truth marginals (LR 53.5%, HR 17.4%, UR 29.1%). Figure 10 shows the per-class deviation ($\Delta\%$, predicted minus ground truth) for each model. Five models showed significant distributional shifts after BH correction (Q range: 23–143, all $p < 0.001$): Gemini 2.5 Pro, Gemini 2.5 Flash, GPT-5.4, Claude Haiku 4.5, and Qwen 3.7-Plus. DeepSeek v4-Flash, GLM-5, and Grok 4.3 showed small, non-significant deviations ($Q \leq 6$, $p \geq 0.04$).

The pattern of deviation was model-specific. Gemini 2.5 Pro strongly overused High Risk (+14.5 pp above the ground-truth rate) while underusing both Low Risk (−7.4 pp) and Unclear Risk (−7.1 pp)—a global shift towards the intermediate-severity class. GPT-5.4 showed the largest Low Risk deficit (−9.4 pp) with compensating Unclear Risk overuse (+7.4 pp), reflecting a tendency to defer to the “insufficient evidence” label. By contrast, Claude Haiku 4.5 and Qwen 3.7-Plus both overused Low Risk markedly (+16.2 and +11.2 pp, respectively) at the expense of Unclear Risk, consistent with the high leniency rates observed in their

531 per-domain error profiles.




Figure 10: Stuart–Maxwell deviation plot. Each model row shows the difference between the predicted class proportion and the ground-truth marginal (Low Risk 53.5%, High Risk 17.4%, Unclear Risk 29.1%) across all 1,589 domain–study cells. Models marked with * differ significantly from the ground-truth distribution after Benjamini–Hochberg FDR correction ($\alpha = 0.05$); the Q statistic (χ^2 , $df = 2$) is shown for each model.

3.2 Performance by RoB Domain

Domain-level accuracy (Figure 11; Table SS2) revealed two tiers of difficulty. **D1** (sequence generation) and **D2** (allocation concealment) were the most accurately assessed domains, with mean accuracy of 75.3% and 72.7%, respectively; all models exceeded 70% accuracy on D1. **D3** (blinding of participants) showed high variance across models (49.8%–72.2%), with Gemini 2.5 Flash and Claude Haiku 4.5 outperforming others by a margin of more than 5 percentage points.

D7 (other sources of bias) was the most challenging domain for all models, with accuracy ranging from 24.2% (Gemini 2.5 Pro) to 55.5% (Qwen 3.7-Plus)—a spread of 31 percentage points within a single domain. Notably, models at opposite ends of the D7 accuracy range adopted diametrically opposed prediction strategies: Gemini 2.5 Pro assigned High Risk to over 93% of D7 cells (recall 0.94 for High Risk; recall 0.18 for Low Risk), whereas Qwen 3.7-Plus assigned Low Risk to 88% of D7 cells (recall 0.88 for Low Risk; recall 0.06 for High Risk). Both strategies yielded similarly poor Macro-F1 in D7 (≤ 0.37 for both), demonstrating that high accuracy on a single class does not translate to overall competence.

The domain-level pattern identifies two structurally distinct types of difficulty. **D4** (blinding of outcome assessors) exhibits systematic, directionally consistent errors: across all eight models, the large majority of errors are lenient—models assign Low Risk when the ground truth is High Risk (see Figure 13)—a pattern present across all intervention types. In drug trials in particular, models appear to conflate the objectivity of the outcome measure with the adequacy of blinding of the outcome assessor, defaulting to Low Risk when no explicit blinding failure is stated. **D7** (other sources of bias), by contrast, is characterised by inter-model divergence: models adopt idiosyncratic strategies—some biased toward High Risk, others toward Low Risk—so that errors are distributed roughly equally between lenient and cautious directions. The wide accuracy spread across models in D7 (31 pp) and its near-zero chance-corrected agreement thus reflect not a shared systematic error but genuine criterion ambiguity; the residual difficulty in D7 is




Figure 11: Per-domain accuracy by model (95% bootstrap CIs; study-clustered, 5,000 iterations, resampling 227 studies with replacement). Each group of eight bars corresponds to one RoB domain; colours identify models as shown in the legend. The dashed horizontal line marks 50% (majority-class baseline). All analyses are exploratory. See Figure 12 for the breakdown by Macro-F1, per-class F1, and Gwet's AC1.

unlikely to be resolved by prompt refinement alone.

Model agreement by domain. Table 5 disaggregates the consensus failure distribution by domain. D1 and D2 are the easiest: 54.2% and 46.3% of their cells achieve full eight-model consensus, and consensus failure rates are 11.0% and 8.8%, respectively. D7 (other sources of bias) stands apart in a different way: only 4.0% of its cells are answered correctly by all models, and 45.8% are answered correctly only by a minority (1–3 models), reflecting a domain where models systematically diverge in their judgements. D4 and D5 present a third pattern: their consensus failure rates (22.0% and 19.8%) are the highest in the corpus—higher even than D7 (12.3%)—indicating that when models err in D4/D5 they tend to do so collectively, whereas in D7 errors are more dispersed across models.

Table 5: Per-cell agreement distribution by domain (row percentages; $n = 227$ per domain).

Domain	All (8/8)	Majority (5–7)	Split (4)	Minority (1–3)	None (0)
	%	%	%	%	%
D1	54.2	22.5	1.3	11.0	11.0
D2	46.3	27.3	2.6	15.0	8.8
D3	33.5	28.2	7.0	17.6	13.7
D4	32.2	16.3	6.6	22.9	22.0
D5	30.4	25.1	7.9	16.7	19.8
D6	19.4	37.4	6.2	25.1	11.9
D7	4.0	27.3	10.6	45.8	12.3

Per-class breakdown. Figure 12 disaggregates performance into the three constituent classes across all domains. Across all models, the ordering $F1\text{-}LR \gg F1\text{-}UR \geq F1\text{-}HR$ holds as a near-universal pattern, closely mirroring the prevalence ordering (53.5%, 29.1%, and 17.4%, respectively; Table 1). The sole exception is Claude Haiku 4.5, which inverts the UR–HR ordering ($F1\text{-}HR = 0.469 >$

578 F1-UR = 0.413): Haiku is more willing to assign High Risk when text signals a
579 methodological concern, at the cost of under-using Unclear Risk as an epistemic
580 hedge—consistent with its highest F1-LR (0.743) and lowest F1-UR (0.413) in the
581 benchmark.

582 **Low Risk** detection was strong and consistent across models in D1 and D2
583 ($F1 > 0.75$ for all models), but deteriorated in D7: most models reported F1-LR
584 below 0.60 there, and Gemini 2.5 Pro collapsed to $F1 = 0.29$ as a consequence of
585 its High Risk-biased prediction strategy.

586 **High Risk** detection was reliable only in D3 ($F1: 0.47\text{--}0.79$), the single domain
587 where High Risk cases outnumbered Low Risk in the ground truth (103 vs. 74;
588 Table 1). The wide bootstrap CIs for F1-HR in D1 and D2 (Figure 12) reflect the small
589 number of High Risk cases in those domains (D1: $n = 5$; D2: $n = 10$), such that a
590 single misclassified study shifts the F1 estimate substantially under resampling.
591 In D6 (selective outcome reporting), four models (Gemini 2.5 Flash, GPT-5.4,
592 Claude Haiku 4.5, and Qwen 3.7-Plus) reported High Risk $F1 = 0.00$, corresponding
593 to zero true-positive High Risk predictions across all 227 studies: High Risk is the
594 least prevalent class in D6 (25 of 227 cases, 11%), and these models systematically
595 defaulted to Low Risk or Unclear Risk in the absence of explicit evidence of
596 selective reporting.

597 **Unclear Risk** was the most consistently under-detected class. Mean recall
598 fell below 0.20 in D3, D5, and D7 across all models. The worst single result
599 was Claude Haiku 4.5 in D5: recall = 0.04 (2 of 45 Unclear Risk cases correctly
600 identified). The most accurately detected Unclear Risk domain was D2, where
601 all models achieved $F1 > 0.70$, consistent with the near-equal prevalence of Low
602 Risk and Unclear Risk in that domain (113 vs. 104). The pattern of under-detection
603 reflects an asymmetric decision behaviour: in D3, D4, D5, and D7—where the
604 ground-truth Unclear Risk rate ranges from 20% to 29%—all models committed
605 to definitive Low Risk or High Risk judgements more often than the Cochrane
606 reviewer, even in cases judged as informationally insufficient. In D2 and D6 the

pattern reverses, with models over-using Unclear Risk relative to the ground truth. This bidirectional discrepancy indicates that models calibrate their uncertainty threshold differently across domains; reducing false positives for Unclear Risk in the blinding domains (D3–D4) is a necessary condition for improving High Risk recall there. One prompt-design factor may contribute to the over-use of Unclear Risk in D6: the prompt specifies a *compound* criterion for that domain (“no protocol/registry **and** insufficient information”), requiring both conditions simultaneously, whereas all other domains require only “insufficient information” to trigger Unclear Risk. This asymmetry may lower the Unclear Risk threshold in D6 relative to other domains, resulting in higher Unclear Risk assignment rates and correspondingly lower Low Risk rates.

Error direction by domain.

Among cells where models erred, errors were not directionally neutral. Figure 13 decomposes errors into three categories under the LR-anchor framework: *lenient* (Low Risk assigned when High Risk or Unclear Risk was correct), *cautious* (High Risk or Unclear Risk assigned when Low Risk was correct), and *lateral* (High Risk and Unclear Risk confused with each other—the study remains flagged but with a mislabelled severity).

D4 (blinding of outcome assessors) exhibited the most pronounced and consistent leniency (50–74% across intervention types), with drug/medication trials showing the highest lenient fraction (74%) and lateral confusion remaining low ($\leq 35\%$). Models appear to conflate the objectivity of the outcome measure with adequate assessor blinding, defaulting to Low Risk when no explicit blinding failure is stated.

D3 (blinding of participants and personnel) showed both high leniency and the highest lateral fraction in the benchmark. While leniency was elevated across all intervention types (48–76%), the lateral fraction was highest for non-pharmacological interventions (42%) and medical devices (24%), indicating that models often recognised a blinding concern without resolving whether it war-

636 ranted High Risk or Unclear Risk—a distinction that turns on reporting quality
637 rather than intervention feasibility.

638 **D1** (sequence generation) and **D2** (allocation concealment) showed a pre-
639 dominantly *cautious* error profile: when models erred on these domains, they
640 over-flagged Low Risk studies as High Risk or Unclear Risk (55–87% cautious
641 in D1; 49–100% in D2). This reversal relative to D3–D4 is consistent with the
642 explicit, formulaic language used to describe randomisation procedures—models
643 rarely miss a clear statement of adequate randomisation and instead raise false
644 concerns when reports are terse.

645 **D6** (selective outcome reporting) was also predominantly cautious (48–72%),
646 reflecting a tendency to flag concerns in the absence of a registered protocol
647 even when the ground-truth judgement was Low Risk.

648 **D7** (other sources of bias) showed no dominant direction, with roughly bal-
649 anced lenient and cautious fractions and a lateral component of 12–21%—confirming
650 the idiosyncratic, criterion-ambiguous character of errors in this domain noted
651 above.



Figure 12: Per-domain performance by model across five metrics: Cohen's κ , F1 for Low Risk, Unclear Risk, and High Risk, and Gwet's AC1 (95% bootstrap CIs; study-clustered, 5,000 iterations, resampling 227 studies with replacement). Each panel shares the same domain grouping on the x -axis; bars are point estimates; error bars are 95% CIs. Short horizontal segments show domain-specific analytical baselines (*Always LR* green, *Always HR* red, *Always UR* orange), computed from each domain's class prevalence; the blue dotted line is the *Random* uniform baseline (constant 1/3 across domains). Baselines with F1 = 0 for a given metric are not drawn. For κ and AC1, the dashed grey line marks zero (chance-level agreement). Macro-F1 by domain is reported in Table SS3. Colours identify models consistently throughout the paper.

Error direction by domain • intervention type (LR-anchor framework)

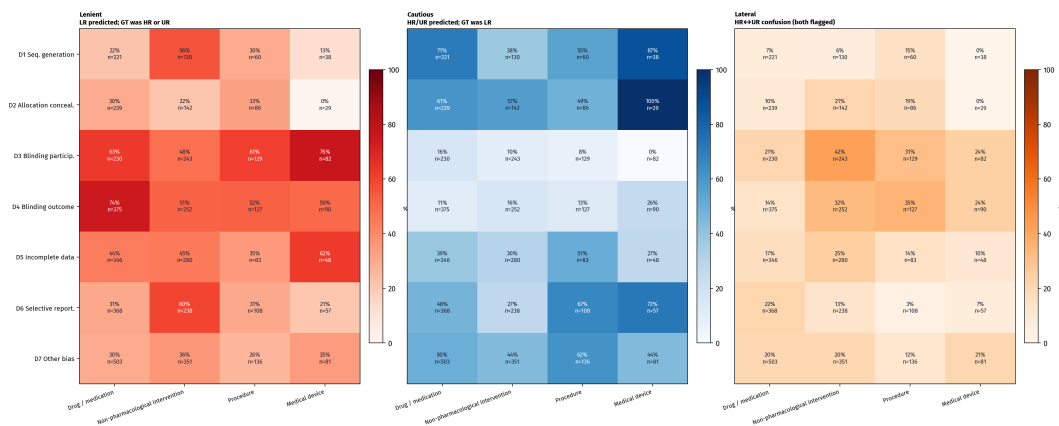


Figure 13: Error direction by RoB domain and intervention type under the LR-anchor framework (three panels, pooled across all eight models). **Lenient** (red): model assigns Low Risk when ground truth was High Risk or Unclear Risk—the study passes unscrutinised. **Cautious** (blue): model assigns High Risk or Unclear Risk when ground truth was Low Risk—a valid study is flagged for unnecessary review. **Lateral** (orange): model confuses High Risk with Unclear Risk or vice versa—the study remains in the flagged pile but with a mislabelled severity. The three percentages sum to 100% in each cell; n = total errors for that domain–intervention combination (identical across all three panels). Grey cells indicate no errors in that combination.

3.3 Stratification by Intervention Type

Domain accuracy by model — stratified by intervention type (95% bootstrap CIs)




Figure 14: Per-domain accuracy by model, stratified by intervention type (95% bootstrap CIs; study-clustered, 5,000 iterations). Panels from top to bottom: drug/medication ($n = 104$), non-pharmacological intervention ($n = 75$), procedure ($n = 30$), and medical device ($n = 18$). Bootstrap CIs for the two smaller strata are wide and should be interpreted with caution; see Supporting Information S4 for numeric estimates.

Mean Macro-F1 across all models ranked intervention types as follows: non-pharmacological intervention (0.572, $n = 75$ studies), drug/medication (0.538, $n = 104$), medical device (0.531, $n = 18$), and procedure (0.510, $n = 30$). The 6.2-point gap between non-pharmacological interventions and procedures was consistent in direction across most models; no single model reversed this ordering. However, bootstrap CIs for the two smaller strata (medical device, $n = 18$; procedure, $n = 30$) are wide (see Figure 14), so this pattern should be interpreted as a descriptive observation rather than a reliable rank ordering pending replication in larger strata. Individual model rankings within each intervention type are shown in Figure 14 and provided in full in Table SS4.

The interaction between intervention type and blinding assessment reveals a specific failure mode. In **D3** (blinding of participants and personnel), accuracy for drug/medication trials substantially exceeded that for device and procedural trials; in the latter, the ground-truth High Risk prevalence is highest—because physical concealment of the intervention is frequently impossible—yet models assigned Low Risk to the majority of cells, producing leniency rates above 65% (Figure 13). In **D4** (blinding of outcome assessors), the leniency bias reached its maximum in drug trials (exceeding 80%), where models appear to assume that objective outcome measurement implies adequate assessor blinding. These findings converge on a specific prompt deficiency: the current zero-shot prompt does not distinguish between trial types when applying blinding criteria. Models apply a “blinding is feasible and likely” default regardless of the physical feasibility of concealment, an error that a single conditional rule in the prompt (e.g., flagging as High Risk by default when concealment is structurally impossible and no blinding mechanism is described) would substantially reduce.

3.4 Stratification by Clinical Domain

Among clinical domains with at least five studies, mean Macro-F1 across models ranged from 0.425 (Infectious Disease, $n = 8$) to 0.680 (Orthopaedics & Trauma,

681 $n = 7$)—a 25.5-point spread (Figure 15). The four highest-performing domains
682 were Orthopaedics & Trauma (0.680), Endocrine & Metabolic (0.638, $n = 10$),
683 Rheumatology (0.637, $n = 5$), and Wounds (0.628, $n = 5$). The four lowest-
684 performing domains were Infectious Disease (0.425), Reproductive & Sexual
685 Health (0.437, $n = 8$), Genetic Disorders (0.473, $n = 9$), and Mental Health (0.481,
686 $n = 11$). Complete results for all 24 Cochrane Review Groups are provided in
687 Table SS5.

688 These results are hypothesis-generating rather than confirmatory. Group sizes
689 of 5–11 studies at the extremes preclude reliable rank-ordering across domains,
690 and multi-label tagging makes groups non-exclusive (some reviews are cross-
691 listed under multiple Cochrane Review Groups). The pattern is nonetheless qual-
692 itatively consistent with an interpretable explanation: domains characterised by
693 standardised reporting (orthopaedic interventions and metabolic trials typically
694 follow structured protocols with explicit randomisation and blinding procedures)
695 show higher Macro-F1, while domains with heterogeneous or complex study
696 designs (Infectious Disease, Reproductive & Sexual Health) present harder RoB
697 characterisation challenges for models trained on general text. These observa-
698 tions motivate targeted future evaluation on larger domain-specific corpora.

699 **3.5 Inference Mode Comparison**

700 Table 6 compares per-domain and single-call inference modes for the two models
701 evaluated under both conditions.

702 The overall picture from this two-model pilot comparison (Table 6) is one of
703 small differences with wide uncertainty: of 12 Δ values across both models and
704 six metrics, ten have 95% CIs that include zero. The two whose CIs exclude zero
705 are a drop in Unclear Risk F1 for Gemini 2.5 Flash (-0.044 ; CI $[-0.079, -0.009]$, †)
706 and a gain in Low Risk F1 for DeepSeek v4-Flash ($+0.026$; $[+0.004, +0.047]$, †); both
707 are class-specific and do not translate into differences in accuracy, Macro-F1,
708 or AC1, all of which remain within their respective uncertainty intervals. The




Figure 15: Macro-F1 by clinical domain and model (domains with ≥ 5 studies; n = number of studies in each domain, multi-label). Colour scale: 0–1.

Table 6: Per-domain vs. single-call inference mode comparison for Gemini 2.5 Flash and DeepSeek v4-Flash ($n = 227$ studies; 1,589 cells each). F1 scores are computed globally across all cells (identical methodology to Table SS1); per-domain values reproduce the corresponding entries in Table SS1. Δ : single-call minus per-domain; 95% paired bootstrap CI (study-clustered, 5,000 iterations) shown in brackets. †: CI excludes zero (unadjusted for multiplicity; all analyses exploratory).

Model	Accuracy (%)	Macro-F1	F1-LR	F1-HR	F1-UR	Gwet’s AC1
<i>Gemini 2.5 Flash</i>						
Per-domain	60.7	0.558	0.718	0.474	0.484	0.437
Single-call	59.8	0.548	0.715	0.490	0.440	0.425
Δ	−0.9	−0.010	−0.003	+0.016	−0.044†	−0.012
	[−2.7, +1.0]	[−0.030, +0.011]	[−0.019, +0.014]	[−0.013, +0.051]	[−0.079, −0.009]	[−0.038, +0.015]
<i>DeepSeek v4-Flash</i>						
Per-domain	60.3	0.558	0.704	0.480	0.489	0.433
Single-call	62.6	0.576	0.730	0.475	0.523	0.466
Δ	+2.3	+0.018	+0.026†	−0.005	+0.034	+0.033
	[−0.1, +4.7]	[−0.007, +0.044]	[+0.004, +0.047]	[−0.049, +0.041]	[−0.003, +0.072]	[−0.001, +0.066]

accuracy differences (−0.9 pp for Flash, +2.3 pp for DeepSeek), while pointing in opposite directions, are each compatible with zero change given the bootstrap CIs ([−2.7, +1.0] and [−0.1, +4.7] respectively).

Domain-level results reinforce this picture. D7 (Other Bias), the domain for which joint assessment was hypothesised to confer the greatest benefit through shared contextual reasoning, showed negligible differences under both modes (± 1 pp)—the most direct evidence that collapsing seven domain assessments into a single call does not meaningfully improve the handling of the hardest domain. The opposing D3 pattern (Flash: −0.4 pp; DeepSeek: −3.6 pp) and the concentration of DeepSeek gains in D5 (+10.2 pp) and D2 (+4.4 pp) suggest that whatever domain-level shifts occur are idiosyncratic to each model rather than reflecting a systematic benefit of joint assessment.

Single-call outputs were systematically shorter (DeepSeek: 400→214 chars; Gemini 2.5 Flash: 530→354 chars), with a corresponding decrease in hedging expressions (e.g., “unclear”, “not reported”, “cannot determine”) in free-text justifications, consistent with a more compressed generation regime under the single-call prompt. The practical implications of this output compression appear limited and model-dependent: it did not translate into a consistent performance benefit or detriment across either model at the aggregate level.

4 Conclusions

Eight state-of-the-art LLMs assessed RoB 1.0 for 227 randomised controlled trials in a zero-shot, single-agent setting, addressing whether the reliability shortfall documented in prior work is uniform across the RoB 1.0 structure or varies predictably by domain and intervention type. Overall accuracy ranged narrowly from 57.6% to 62.5%, but Macro-F1 and Gwet's AC1 revealed meaningful differences in how models handled the three nominal classes. Models with the highest accuracy achieved it partly by concentrating predictions on majority-class labels, a strategy that inflates accuracy whilst suppressing detection of High Risk and Unclear Risk judgements—the very categories most consequential for evidence appraisal. Macro-F1 and Gwet's AC1 are therefore more informative primary metrics than accuracy for this task.

Unclear Risk was the most consistently under-detected class across all models and domains, with recall below 0.20 in three of seven domains. This systematic failure reflects the epistemic character of the category—it requires the model to recognise that key methodological information is absent, a more complex inference than identifying a positive or negative methodological signal in the text.

Domain-level and stratification analyses revealed structured heterogeneity in performance. Sequence generation (D1) and allocation concealment (D2) were reliably assessed across all models, whereas other sources of bias (D7) showed the greatest inter-model accuracy spread and blinding of outcome assessors (D4) had the highest consensus failure rate (22.0%—no model correct in 50 of 227 cells). Clinical domains with clear, domain-specific methodological criteria (Orthopaedics & Trauma, Endocrine & Metabolic) supported higher Macro-F1 than domains with more heterogeneous study designs (Infectious Disease, Reproductive & Sexual Health).

Consensus failures—cells where no model was correct—accounted for 14.2% (226 of 1,589 cells) and were concentrated in D4 and D5, the two most ambigu-

ous domains. These cases identify the boundaries of current frontier model capability under this evaluation design and may reflect genuine annotation uncertainty as much as model failure. Pairwise complementarity analysis showed that 7.4%–17.4% of cells are correctable by substituting any one model for another, confirming that the eight models make substantially non-redundant errors.

The comparison between per-domain and single-call inference modes (Section 3.5) produced small differences that do not collectively support a performance advantage for either strategy. Ten of twelve Δ values across both models and six metrics had 95% CIs spanning zero; the two exceptions were class-specific effects of limited practical magnitude (Unclear Risk F1 for Flash: -0.044 ; Low Risk F1 for DeepSeek: $+0.026$) that did not propagate to aggregate metrics. Most notably, D7 (Other Bias)—the domain for which joint contextual assessment was hypothesised to be most beneficial—showed negligible differences under both modes (± 1 pp), indicating that its difficulty is intrinsic to the vagueness of the criterion rather than a consequence of fragmented inference calls. Taken together, the preliminary evidence from this two-model comparison does not identify a meaningful performance rationale for preferring single-call over per-domain assessment; the primary practical distinction between the two modes is computational (single-call reduces inference calls by $7\times$) rather than one of accuracy.

Reframing High Risk detection as a binary classification task (High Risk vs. not High Risk) highlights the clinical stakes of the leniency bias documented above. Under this framing, the best-performing models for High Risk recall detected fewer than two thirds of High Risk cells (F1-HR: 0.455 – 0.480 for the top-ranked models), meaning roughly one in three cells that should trigger reviewer scrutiny passed undetected. Both error directions carry clinical costs: a false negative (lenient—biased study classified as acceptable and included unquestioned in a meta-analysis) risks inflating pooled effect estimates, while a false positive (cautious—a low-risk study flagged as problematic) risks unjustifiably discrediting

786 or excluding credible evidence, a concern that grows when the pool of eligible
787 studies is small. The relative severity depends on review context; however,
788 tools intended to assist real-world screening workflows would need to make
789 this trade-off explicit rather than leaving it to the zero-shot prompt's implicit
790 defaults.

791 The cell-level concordance partition offers a practical framework for directing
792 qualitative review effort. Cells in which all models converge on the correct
793 answer require no additional attention. Cells in which all models converge on
794 the *same wrong* answer are the highest-priority targets for prompt refinement
795 or annotation review: unanimity of error across architecturally diverse models
796 makes model-idiosyncratic explanation implausible, pointing instead to shared
797 training priors, prompt ambiguity, or genuine ground-truth noise. Cells in which a
798 single model is the sole outlier—either the only model correct or the only model
799 incorrect—index model-specific capacity gaps independent of task difficulty.
800 Targeting qualitative analysis at consensus errors before individual outliers
801 maximises diagnostic yield per unit of human effort.

802 The domain and stratification analyses converge on a specific, actionable
803 direction for prompt improvement: a blinding rule sensitive to intervention
804 type. Models would benefit from explicit guidance distinguishing trials in which
805 blinding is physically feasible (pharmacological interventions with matched
806 placebos) from those in which it is not (surgical procedures, physiotherapy,
807 medical devices), with the latter receiving a conservative default that does not
808 presuppose the availability of a blinding mechanism. This single rule targets
809 D4—the domain with the most systematic and directionally consistent errors
810 in the benchmark—and has direct implications for the High Risk recall deficit
811 identified above.

812 The recent preregistered evaluation by Nyrhi et al. [9] demonstrates that
813 automated RoB assessment with frontier LLMs is an active area of investigation.
814 That study applied four models (ChatGPT o3, DeepSeek v3, Gemini Flash 2.0,

Grok 3) to 100 RCTs from 10 Cochrane reviews across four clinical specialties, using multimodal PDF input, and reported interobserver κ of 0.27–0.39 for RoB 1.0. The present study and Nyrhi et al. share no overlapping models and differ substantially in corpus composition (127 vs. 10 Cochrane reviews; 24 vs. 4 clinical specialties), corpus size (227 vs. 100 RCTs), and input modality (text-only pipeline vs. direct PDF upload). Direct numerical comparison of performance figures is therefore not warranted. Qualitatively, the κ values observed here ($\kappa = 0.325$ for the best model; lower for the remaining seven) fall within and below the range they reported, which is consistent with expectations given the broader and more heterogeneous corpus used here—more diverse corpora generally present harder classification problems—and the use of a text-only rather than multimodal pipeline. Both studies independently reach the same qualitative conclusion: contemporary frontier LLMs do not meet the reliability threshold for autonomous RoB assessment under a zero-shot, unassisted design.

The present study contributes three elements that go beyond that aggregate conclusion. First, evaluating eight models spanning five provider families on a corpus of 227 studies across 24 clinical specialties reveals that aggregate performance is approximately stable across frontier models, shifting the central question from *which model* to *which domains and intervention types* support reliable assistance—a question that aggregate summaries cannot answer. Second, domain-stratified analysis identifies structured heterogeneity: D1 and D2 may represent nearer-term targets for assisted screening, whereas the systematic leniency bias in D4 and the idiosyncratic divergence in D7 identify specific failure modes that require targeted prompt improvements before reliable assistance is achievable. Third, explicit degenerate baselines and the juxtaposition of Cohen’s κ with Gwet’s AC1 expose analytical pitfalls specific to this task: the margin above trivially achievable strategies is narrowest for the clinically consequential classes (High Risk and Unclear Risk), and the two agreement coefficients diverge most sharply for the model with the highest accuracy (Claude Haiku 4.5: $\kappa = 0.325$

844 vs. $AC1 = 0.485$), cautioning against treating either coefficient in isolation when
845 class distributions are skewed.

846 Several limitations should be noted. First, this study was not pre-registered;
847 the analysis plan was developed alongside the benchmark construction, and
848 the choice of metrics, stratification variables, and comparisons was informed
849 by exploratory review of the data. Findings should therefore be interpreted as
850 exploratory and hypothesis-generating throughout, notwithstanding the use
851 of formal confidence intervals to characterise estimation uncertainty; they are
852 not the outcome of pre-specified confirmatory hypothesis tests. Second, the
853 ground-truth annotations derive from published Cochrane review reports and
854 were not independently re-annotated for this study. The Cochrane annotations
855 represent the consensus output of a structured multi-reviewer process (independ-
856 ent assessment by at least two trained reviewers, with disagreements resolved
857 by discussion or arbitration), which is the reference standard used operationally
858 by the field and the most appropriate benchmark for evaluating a tool intended
859 to assist real-world systematic review workflows. Quantifying internal reliability
860 within this corpus would require re-annotating a subsample with independent
861 trained reviewers; however, agreement between Cochrane-trained reviewers
862 and non-specialist re-annotators would conflate task difficulty with expertise
863 asymmetry—the same confound that limits the interpretability of comparisons
864 between Cochrane and external reviewers [5]. The relevant population-level IRR
865 estimate is therefore taken from Hartling et al. [1], who measured agreement
866 between formally trained reviewer pairs applying the same protocol: $\kappa = 0.79$ for
867 D1 but only $\kappa = 0.24$ – 0.37 for remaining domains, falling to $\kappa = 0.05$ – 0.09 between
868 independent consensus pairs; independent Cochrane teams show similarly vari-
869 able agreement [4]. These figures indicate that the reference standard itself
870 carries substantial noise, placing a practical ceiling on any automated system’s
871 agreement with published annotations that is independent of model quality.
872 The observed model accuracy range (57.6%–62.5%) therefore reflects agreement

with a reference standard that is itself imperfect across most non-D1 domains: disentangling genuine model failure from annotation noise in the ground truth is not possible without independent re-annotation of the benchmark corpus. Third, each model was queried once per study using provider-default generation parameters. This reflects the intended evaluation scope: zero-shot performance as models are available to practitioners, without artificial parameter tuning. Manipulating generation temperature would introduce heterogeneous conditions across model families, and strategies such as majority voting over repeated draws represent post-hoc optimisation outside the zero-shot benchmark design. As a consequence, differences between models separated by fewer than five percentage points should not be interpreted as definitive rank orderings in the absence of confidence intervals; these are discussed as directional trends. Fourth, all models were accessed in June 2026 via the OpenRouter API using the identifier strings reported in Table 2; providers may silently update the weights associated with a fixed identifier, so exact reproduction requires accessing the same model strings within the same inference window. Results should therefore be interpreted as reflecting the checkpoint state active in June 2026, not any future revision of the same identifier. Raw model outputs (per-cell predictions and justifications) are not publicly deposited: the full-text PDFs used as input are not all open-access and re-distribution could infringe publisher copyright; additionally, the underlying study corpus is currently under independent expert medical review for a forthcoming benchmark publication, and its public release is deferred to that work. Fifth, the single-agent zero-shot design examines one specific slice of the capability space; few-shot prompting, retrieval augmentation, fine-tuning, and multi-agent architectures were not evaluated and may yield substantially different results. Sixth, training-data contamination cannot be excluded, but several structural features of the evaluation design limit its plausibility as a primary driver of results. Memorisation-based performance would require a model to (i) have seen the particular one of the 127 distinct

902 Cochrane 2024 reviews from which each study was drawn, (ii) correctly iden-
903 tify that the input Markdown document corresponds to a specific entry in that
904 review’s RoB table, and (iii) reproduce the reviewer’s label rather than reason
905 from the text. The Cochrane reviews sampled here were published in 2024;
906 most included RCTs appear in paywalled journals, making full-text ingestion
907 at training time less likely for the majority of papers. Moreover, the empirical
908 pattern is inconsistent with memorisation at scale: the best individual model
909 (Claude Haiku 4.5) achieved a perfect 7/7 score on only 12 of 227 studies (5.3%),
910 and even across the full ensemble of eight models all seven domains were correct
911 in only 40 studies (17.6%). Both figures are consistent with the rates expected
912 under independent domain judgements at the observed per-domain accuracy
913 ($0.60^7 \approx 2.8\%$; inflated modestly by the low intraclass correlation of $\hat{\rho} = 0.044$ to
914 roughly 4–5% per model; see Supporting Information S7 for derivation), leaving
915 no systematic excess attributable to retrieval. Targeted contamination probes
916 (e.g. membership inference or verbatim reproduction tests) were not conducted,
917 as the opaque training corpora of the evaluated models preclude rigorous anal-
918 ysis; this remains a caveat for interpretation. Seventh, the Cochrane corpus
919 used here comprises better-reported trials than the broader trial literature, so
920 these performance estimates may be optimistic relative to unselected primary
921 research.

922 **Acknowledgements**

923 We gratefully acknowledge the financial support of the TELUS Digital Research
924 Hub.

925 **Data and Code Availability**

926 The analysis code is publicly available at <https://github.com/ACS-USP/systematic-reviews>.

Conflicts of Interest

The authors declare no conflicts of interest.

References

- [1] Lisa Hartling, Michele P Hamm, Andrea Milne, Ben Vandermeer, P Lina Santaguida, Mohammed Ansari, Alexander Tsertsvadze, Susanne Hempel, Paul Shekelle, and Donna M Dryden. Testing the risk of bias tool showed low reliability between individual reviewers and across consensus assessments of reviewer pairs. *Journal of Clinical Epidemiology*, 66(9):973–981, 2013. doi: 10.1016/j.jclinepi.2012.07.005.
- [2] Julian PT Higgins, James Thomas, Jacqueline Chandler, Miranda Cumpston, Tianjing Li, Matthew J Page, and Vivian A Welch, editors. *Cochrane Handbook for Systematic Reviews of Interventions*. John Wiley & Sons, 2nd edition, 2019. doi: 10.1002/9781119536604.
- [3] Julian PT Higgins, Douglas G Altman, and Jonathan AC Sterne. Assessing risk of bias in included studies. In Julian PT Higgins and Sally Green, editors, *Cochrane Handbook for Systematic Reviews of Interventions*, chapter 8. The Cochrane Collaboration, 5.1.0 edition, 2011.
- [4] Nadja Könsgen, Ognjen Barcot, Simone Heß, Livia Puljak, Käthe Goossen, Tanja Rombey, and Dawid Pieper. Inter-review agreement of risk-of-bias judgments varied in Cochrane reviews. *Journal of Clinical Epidemiology*, 120: 25–32, 2020. doi: 10.1016/j.jclinepi.2019.12.016.
- [5] Susan Armijo-Olivo, Maria Ospina, Bruno R da Costa, Matthias Egger, Humam Saltaji, Jorge Fuentes, Christine Ha, and Greta G Cummings. Poor reliability between Cochrane reviewers and blinded external reviewers when applying

- the Cochrane risk of bias tool in physical therapy trials. *PLOS ONE*, 9(5): e96920, 2014. doi: 10.1371/journal.pone.0096920.
- [6] Iain J Marshall, Joël Kuiper, and Byron C Wallace. RobotReviewer: evaluation of a system for automatically assessing bias in clinical trials. *Journal of the American Medical Informatics Association*, 23(1):193–201, 2016. doi: 10.1093/jamia/ocv044.
- [7] Judith-Lisa Lieberum, Markus Toews, Maria-Inti Metzendorf, Felix Heilmeyer, Waldemar Siemens, Christian Haverkamp, Daniel Böhringer, Joerg J Meerpohl, and Angelika Eisele-Metzger. Large language models for conducting systematic reviews: on the rise, but not yet ready for use—a scoping review. *Journal of Clinical Epidemiology*, 181:111746, 2025. doi: 10.1016/j.jclinepi.2025.111746.
- [8] Dmitry Scherbakov, Nina Hubig, Vinita Jansari, Alexander Bakumenko, and Leslie A Lenert. The emergence of large language models as tools in literature reviews: a large language model-assisted systematic review. *Journal of the American Medical Informatics Association*, 32(6):1071–1086, 2025. doi: 10.1093/jamia/ocaf063.
- [9] Lauri Nyrhi, Ville Ponkilainen, Juho Laaksonen, Lauri Kuikka, Lauri Paljakka, Teemu Karjalainen, Ville M Mattila, and Ilari Kuitunen. Large language models for risk-of-bias assessment in randomised clinical trials—a comparative validation study. *eBioMedicine*, 126:106238, 2026. doi: 10.1016/j.ebiom.2026.106238.
- [10] Jonathan AC Sterne, Jelena Savović, Matthew J Page, Pankaj Roy, Julian PT Higgins, Isabelle Boutron, Christopher J Cates, Hao-Yun Cheng, Matthew S Corbett, Sandra M Eldridge, Jonathan R Emberson, Miguel A Hernán, Sally Hopewell, Asbjørn Hróbjartsson, Daniela R Junqueira, Peter Jüni, Jamie J Kirkham, Toby Lasserson, Tianjing Li, Alexandra McAleenan, Barnaby C Reeves, Sasha Shepperd, Ian Shrier, Lesley A Stewart, Kate Tilling, Ian R White,

- 978 Penny F Whiting, and Julian PT Higgins. RoB 2: a revised tool for assessing
979 risk of bias in randomised trials. *BMJ*, 366:l4898, 2019. doi: 10.1136/bmj.l4898.
- 980 [11] Bin Wang, Chao Xu, Xiaomeng Zhao, Linke Ouyang, Fan Wu, Zhiyuan Zhao,
981 Rui Xu, Kaiwen Liu, Yuan Qu, Fukai Shang, Bo Zhang, Liquan Wei, Zhihao Sui,
982 Wei Li, Botian Shi, Yu Qiao, Dahua Lin, and Conghui He. MinerU: An open-
983 source solution for precise document content extraction. *arXiv preprint*
984 *arXiv:2409.18839*, 2024. doi: 10.48550/arXiv.2409.18839.
- 985 [12] Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc V Le, Ed H Chi, Sharan
986 Narang, Aakanksha Chowdhery, and Denny Zhou. Self-consistency improves
987 chain of thought reasoning in language models. In *International Conference*
988 *on Learning Representations*, 2023. URL [https://openreview.net/forum?](https://openreview.net/forum?id=1PL1NIMMrw)
989 [id=1PL1NIMMrw](https://openreview.net/forum?id=1PL1NIMMrw).
- 990 [13] Kilem L Gwet. Computing inter-rater reliability and its variance in the pres-
991 ence of high agreement. *British Journal of Mathematical and Statistical*
992 *Psychology*, 61(1):29–48, 2008. doi: 10.1348/000711006X126600.
- 993 [14] Alvan R Feinstein and Domenic V Cicchetti. High agreement but low kappa:
994 I. The problems of two paradoxes. *Journal of Clinical Epidemiology*, 43(6):
995 543–549, 1990. doi: 10.1016/0895-4356(90)90158-L.
- 996 [15] Quinn McNemar. Note on the sampling error of the difference between
997 correlated proportions or percentages. *Psychometrika*, 12(2):153–157, 1947.
998 doi: 10.1007/BF02295996.
- 999 [16] Alan Stuart. A test for homogeneity of the marginal distributions in a two-
1000 way classification. *Biometrika*, 42(3–4):412–416, 1955. doi: 10.1093/biomet/42.
1001 3-4.412.
- 1002 [17] A. E. Maxwell. Comparing the classification of subjects by two independent
1003 judges. *The British Journal of Psychiatry*, 116(535):651–655, 1970. doi: 10.1192/
1004 bjp.116.535.651.

1005 [18] Yoav Benjamini and Yosef Hochberg. Controlling the false discovery rate: a
1006 practical and powerful approach to multiple testing. *Journal of the Royal*
1007 *Statistical Society: Series B (Methodological)*, 57(1):289–300, 1995. doi: 10.
1008 1111/j.2517-6161.1995.tb02031.x.

1010 **S1 Prompt Used for Single-Agent RoB Assessment**

1011 The following prompt was used without modification across all eight models and
1012 all 227 studies. Each study was assessed in seven sequential inference calls—one
1013 per RoB domain—with the full Markdown text of the study report included in
1014 every call as the user message, followed by the domain identifier (e.g. Domain:
1015 D1). The prompt text below was supplied as the system message.

1016 Goal: Systematically assess the risk of bias in a research article using
1017 Cochrane RoB 1.0. Evaluate bias domain, assign a judgment (Low, High,
1018 or Unclear Risk), and justify strictly from the article evidence.

1019

1020 ---

1021

1022 Preliminary Considerations

1023

1024 Before assessing, identify:

1025 1. Intervention Type: Pharmacological | Surgical/Procedural |

1026 Non-pharmacological | Device

1027 2. Primary Outcome Type: Objective (mortality, lab values) |

1028 Subjective (pain, QoL)

1029 3. Comparator: Placebo | Active comparator | Usual care | No intervention

1030

1031 Critical Context: Blinding expectations vary by intervention type.

1032 For pharmacological interventions, lack of placebo is a serious flaw.

1033 For surgical/non-pharmacological interventions where blinding is impossible,

1034 focus on compensatory measures: blinding of outcome assessors and objective
1035 endpoints.

1036

1037 ---

1038

1039 Bias Domains

1040

1041 1. Random Sequence Generation (Selection Bias)

1042 Signalling Question: Was a truly random method used?

1043 Low Risk: Random method adequately described (computer-generated,
1044 random tables, coin toss, dice, shuffled envelopes,
1045 minimization with random element)

1046 High Risk: Non-random method (alternation, date of birth, hospital
1047 record number, clinician judgment)

1048 Unclear Risk: States "randomized" but no method details provided

1049

1050 2. Allocation Concealment (Selection Bias)

1051 Signalling Question: Was the allocation sequence concealed from those
1052 enrolling participants until assignment?

1053 Low Risk: Adequate concealment (central/telephone randomization,
1054 sequentially numbered opaque sealed envelopes, numbered
1055 drug containers, independent pharmacy)

1056 High Risk: Allocation not concealed or predictable (open list,
1057 non-opaque envelopes, alternation)

1058 Unclear Risk: Insufficient information to judge

1059 Note: Allocation concealment (before randomization) is different from
1060 blinding (after randomization).

1061

1062 3. Blinding of Participants and Personnel (Performance Bias)

1063 Signalling Questions: Were participants/personnel aware of assignments?

1064 If yes, could this affect outcomes?

1065 Pharmacological: Double-blinding expected; no placebo implies higher
1066 risk for subjective outcomes.

1067 Surgical/Non-pharmacological: Blinding may be impossible and is not
1068 automatically High Risk; check compensatory safeguards.

1069 Objective outcomes: Less susceptible to performance bias.

1070 Low Risk: Adequate blinding, OR blinding impossible with compensatory

1071 measures, OR outcome is objective/hard endpoint

1072 High Risk: No blinding + subjective outcome + no compensatory measures

1073 Unclear Risk: Insufficient blinding information

1074

1075 4. Blinding of Outcome Assessors (Detection Bias)

1076 Signalling Questions: Were assessors blinded? Is outcome objective or

1077 subjective?

1078 Low Risk: Assessors blinded, OR objective outcome unlikely to be

1079 influenced by assignment knowledge

1080 High Risk: Assessors not blinded + subjective outcome requiring

1081 interpretation

1082 Unclear Risk: Insufficient information

1083 Note: For patient-reported outcomes, participant acts as assessor.

1084

1085 5. Incomplete Outcome Data (Attrition Bias)

1086 Signalling Questions: Are data available for most participants? Are

1087 missing-data reasons balanced/unrelated to outcome?

1088 Low Risk: Minimal missingness (<5%), OR balanced missingness with

1089 unrelated reasons, OR valid ITT/imputation

1090 High Risk: High missingness (>10-20%), OR imbalance between groups,

1091 OR outcome-related missingness, OR per-protocol exclusions

1092 Unclear Risk: Missingness not adequately reported

1093

1094 6. Selective Reporting (Reporting Bias)

1095 Signalling Questions: Is protocol/registry available? Were pre-specified

1096 outcomes reported?

1097 Low Risk: All pre-specified outcomes reported, OR protocol/registry

1098 aligns with publication

1099 High Risk: Selective non-reporting (methods outcomes missing in

1100 results, changed primary outcome, significance-driven

1101 reporting)

1102 Unclear Risk: No protocol/registry and insufficient information

1103

1104 7. Other Bias

1105 Signalling Questions: Are there other bias sources (early stopping,
1106 baseline imbalance, conflicts, contamination)?

1107 Low Risk: No relevant additional bias detected

1108 High Risk: Important additional bias present

1109 Unclear Risk: Insufficient information

1110

1111 ---

1112

1113 Key Principles

1114 1. Evidence-based only: Base judgments on explicit information. Do not
1115 assume. If evidence is missing, choose "Unclear Risk".

1116 2. Cite evidence: Reference specific text or note absence of information.

1117 3. Context matters: Apply blinding expectations appropriate to intervention
1118 type.

1119 4. Impossibility vs Negligence: Not blinding when impossible differs from
1120 not blinding when feasible.

1121 5. Outcome sensitivity: Subjective outcomes require stricter blinding than
1122 objective ones.

1123

1124 ---

1125

1126 Return ONLY valid JSON with:

1127 {

1128 "judgement": "Low Risk" | "High Risk" | "Unclear Risk",

1129 "justification": "short justification"

1130 }

Table S1: Overall performance on RoB 1.0 assessment ($n = 227$ studies; 1,589 cells). Models ranked by accuracy. F1-LR, F1-HR, F1-UR: F1 for Low Risk, High Risk, and Unclear Risk computed globally across all study-domain cells. Cohen’s κ and Gwet’s AC1 are both chance-corrected agreement coefficients; they differ in how the expected chance agreement p_e is computed (see Methods). Best value in each column shown in bold. 95% confidence intervals (study-clustered bootstrap, 5,000 iterations, resampling 227 studies with replacement) are shown in brackets below each point estimate; all analyses are exploratory (see Methods, Statistical Inference).

Model	Accuracy (%)	Macro-F1	F1-LR	F1-HR	F1-UR	Cohen’s κ	Gwet’s AC1
Claude Haiku 4.5	62.5 [59.8, 65.2]	0.541 [0.511, 0.569]	0.743 [0.718, 0.767]	0.469 [0.412, 0.521]	0.413 [0.361, 0.463]	0.325 [0.286, 0.365]	0.485 [0.443, 0.526]
GLM-5	61.1 [58.7, 63.5]	0.557 [0.529, 0.582]	0.709 [0.682, 0.735]	0.441 [0.385, 0.494]	0.520 [0.479, 0.561]	0.341 [0.301, 0.380]	0.449 [0.411, 0.485]
Qwen 3.7-Plus	60.9 [58.3, 63.4]	0.525 [0.496, 0.552]	0.718 [0.691, 0.743]	0.356 [0.293, 0.413]	0.503 [0.459, 0.543]	0.306 [0.266, 0.345]	0.457 [0.417, 0.495]
Gemini 2.5 Flash	60.7 [58.2, 63.1]	0.558 [0.533, 0.582]	0.718 [0.690, 0.744]	0.474 [0.424, 0.516]	0.484 [0.443, 0.526]	0.349 [0.312, 0.384]	0.437 [0.399, 0.475]
DeepSeek v4-Flash	60.3 [57.7, 62.9]	0.558 [0.531, 0.585]	0.704 [0.676, 0.730]	0.480 [0.429, 0.530]	0.489 [0.447, 0.530]	0.337 [0.296, 0.377]	0.433 [0.394, 0.472]
Grok 4.3	60.1 [57.8, 62.5]	0.549 [0.523, 0.574]	0.696 [0.670, 0.722]	0.432 [0.376, 0.486]	0.518 [0.478, 0.556]	0.326 [0.288, 0.362]	0.433 [0.397, 0.471]
GPT-5.4	57.7 [55.1, 60.3]	0.546 [0.519, 0.572]	0.668 [0.636, 0.698]	0.464 [0.410, 0.511]	0.506 [0.468, 0.543]	0.322 [0.282, 0.361]	0.387 [0.349, 0.426]
Gemini 2.5 Pro	57.6 [55.1, 60.0]	0.543 [0.517, 0.569]	0.674 [0.643, 0.702]	0.455 [0.407, 0.501]	0.502 [0.458, 0.542]	0.331 [0.293, 0.367]	0.383 [0.344, 0.420]

1131 **S2 Overall Model Performance**

1132 **S3 Per-Domain Performance by Model**

1133 Tables S2 and S3 report accuracy (%) and Macro-F1 for each model in each of the
1134 seven RoB domains.

1135 **S4 Stratification by Intervention Type**

1136 Figure S1 shows the distribution of intervention types within each clinical domain
1137 in the 227-study corpus. Because each study can be tagged with multiple clinical
1138 domains, the counts across domains sum to more than 227; a study tagged
1139 with two domains contributes its intervention type to both. The distribution
1140 is markedly non-uniform: Pregnancy & childbirth and Complementary & alternative
1141 medicine are almost exclusively pharmacological, Effective practice & health systems
1142 and Mental health are predominantly non-pharmacological, and
1143 Orthopaedics & trauma contains no pharmacological studies. Table S4 reports
1144 Macro-F1 for each model within each intervention type. Figure S2 shows the same
1145 data as a heatmap.

Table S2: Per-domain accuracy (%) for all eight models ($n = 227$ cells per domain). Best value per domain in bold.

Model	D1	D2	D3	D4	D5	D6	D7
Claude Haiku 4.5	73.6	74.0	67.4	51.5	64.3	57.7	48.9
GLM-5	76.7	74.0	55.9	53.7	62.1	55.9	49.3
Qwen 3.7-Plus	77.5	74.9	49.8	52.0	59.9	56.4	55.5
Gemini 2.5 Flash	72.2	74.0	72.2	51.5	60.4	62.6	31.7
DeepSeek v4-Flash	74.0	70.0	67.0	57.7	54.6	55.9	42.7
Grok 4.3	75.8	70.0	59.5	54.6	58.6	60.8	41.4
GPT-5.4	75.8	70.0	64.3	52.0	51.5	55.9	34.4
Gemini 2.5 Pro	76.7	74.4	62.6	55.1	55.1	55.1	24.2

Table S3: Per-domain Macro-F1 for all eight models ($n = 227$ cells per domain). Best value per domain in bold.

Model	D1	D2	D3	D4	D5	D6	D7
Claude Haiku 4.5	0.572	0.618	0.527	0.425	0.394	0.387	0.269
GLM-5	0.632	0.657	0.462	0.479	0.514	0.443	0.345
Qwen 3.7-Plus	0.638	0.672	0.428	0.447	0.439	0.388	0.288
Gemini 2.5 Flash	0.549	0.663	0.626	0.411	0.412	0.438	0.280
DeepSeek v4-Flash	0.602	0.610	0.583	0.521	0.419	0.424	0.347
Grok 4.3	0.618	0.668	0.524	0.492	0.423	0.449	0.315
GPT-5.4	0.625	0.635	0.515	0.469	0.419	0.397	0.317
Gemini 2.5 Pro	0.621	0.671	0.481	0.478	0.413	0.475	0.195

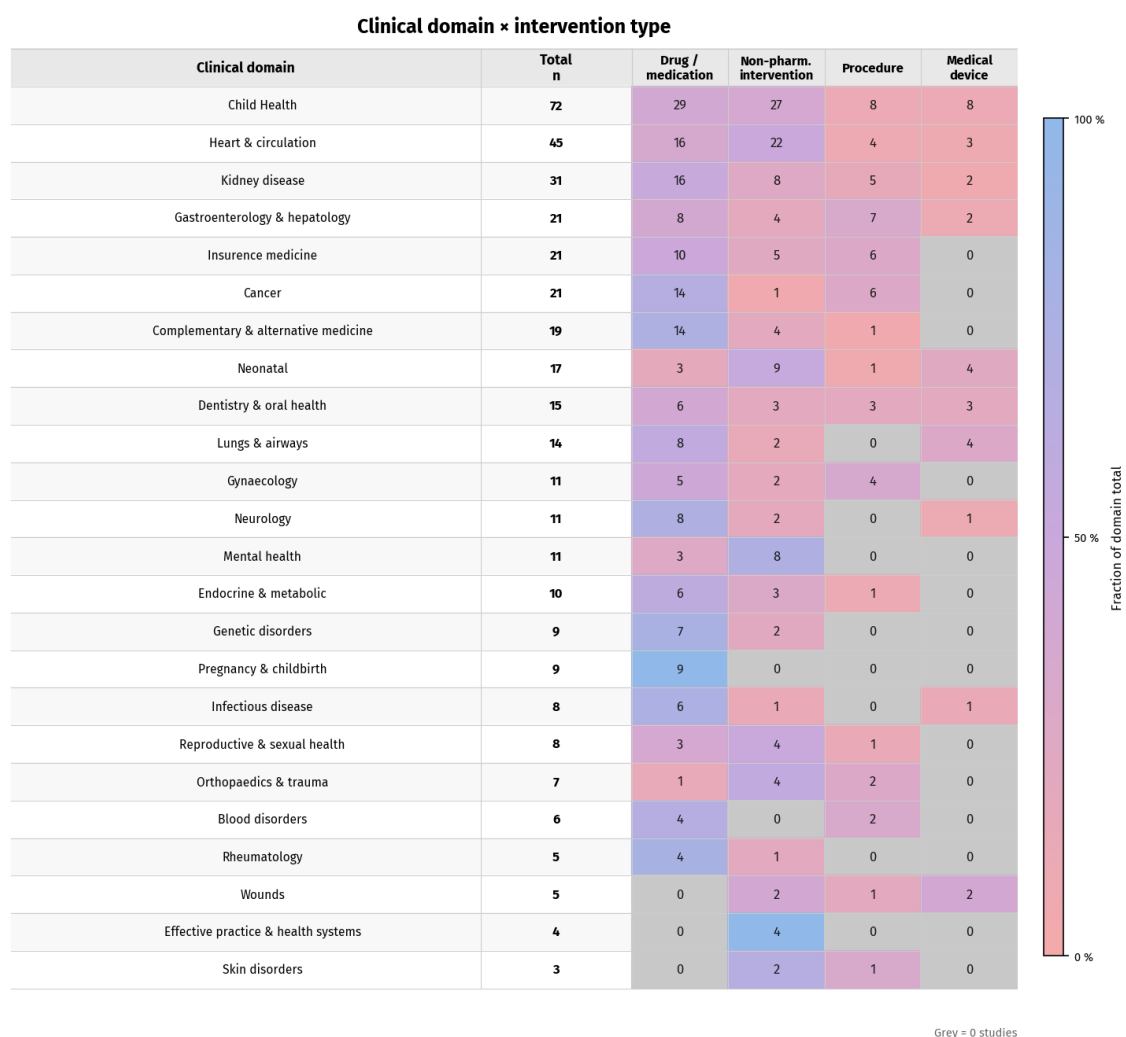


Figure S1: Distribution of intervention types across clinical domains ($n = 227$ studies; studies with multiple clinical domain tags are counted in each relevant domain). Cell colour encodes the fraction of that domain's studies belonging to each intervention type (pastel red = low proportion, pastel blue = high proportion); grey cells indicate zero studies. Domains are sorted by total study count (descending).

S5 Stratification by Clinical Domain

Table S5 reports mean Macro-F1 (averaged across all eight models) for all 24 Cochrane Review Group categories, ordered from highest to lowest mean F1. Domains with fewer than five studies are included for completeness but should be interpreted with caution.

Table S4: Macro-F1 stratified by intervention type. Numbers in parentheses are study counts (n_{studies}). Best value per row in bold. Subgroup estimates carry substantial uncertainty given the small group sizes; these analyses are descriptive only.

Model	Drug / medication ($n = 104$)	Non-pharma- cological ($n = 75$)	Procedure ($n = 30$)	Medical device ($n = 18$)
Claude Haiku 4.5	0.548	0.530	0.494	0.627
GLM-5	0.559	0.590	0.514	0.471
Qwen 3.7-Plus	0.511	0.567	0.474	0.445
Gemini 2.5 Flash	0.540	0.574	0.529	0.596
DeepSeek v4-Flash	0.548	0.582	0.520	0.543
Grok 4.3	0.552	0.561	0.519	0.531
GPT-5.4	0.515	0.595	0.530	0.502
Gemini 2.5 Pro	0.533	0.575	0.498	0.534

S6 Pairwise Model Complementarity

Table S6 reports pairwise complementarity across the eight models. Each cell (i, j) gives the percentage of the 1,589 study–domain cells in which model i produced an incorrect judgement and model j produced the correct one. High values indicate that model j recovers errors that model i misses, i.e. that j adds non-redundant information over i . The matrix is asymmetric: cell $(i, j) \neq$ cell (j, i) because the two models may have different overall accuracy.

S7 Confidence Interval Method Comparison

All confidence intervals reported in the main text were computed using study-clustered bootstrap (5,000 iterations, resampling the 227 studies with replacement). This section documents the comparison between bootstrap and two Wilson score interval variants for accuracy, and explains why the bootstrap approach was adopted for all metrics.

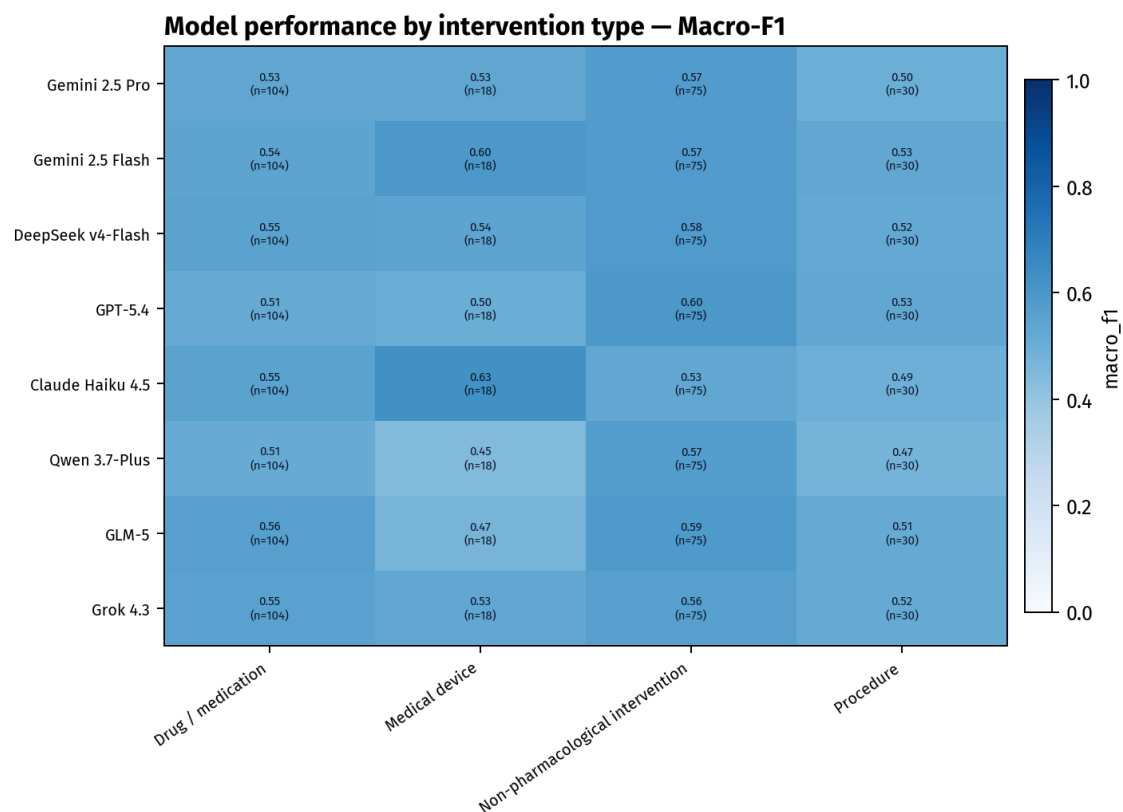


Figure S2: Macro-F1 heatmap by intervention type and model. Colour scale: 0–1.

Wilson score interval

For a proportion p observed over n independent Bernoulli trials, the Wilson 95% CI is

$$\left(\frac{np + \frac{z^2}{2}}{n + z^2} \pm \frac{z \sqrt{np(1-p) + \frac{z^2}{4}}}{n + z^2} \right), \quad (\text{S1})$$

with $z = 1.96$. Applied to accuracy ($p \approx 0.60$), two natural choices for n exist: $n = 1,589$ (all annotated cells, assuming independence) or $n = 227$ (studies as the primary sampling unit).

Intraclass correlation and effective sample size

The 1,589 cells are not independent: seven domain judgements are drawn from the same study. Pearson's one-way intraclass correlation coefficient (ICC) for the correct/wrong binary indicator across the seven domains within each study was estimated at $\hat{\rho} = 0.044$ (computed for Claude Haiku 4.5 as a representative

Table S5: Mean Macro-F1 across all eight models, stratified by clinical domain (24 Cochrane Review Groups, multi-label; study counts sum to more than 227). Ordered by descending mean F1. Estimates for smaller groups carry substantial uncertainty; these analyses are descriptive only.

Clinical domain	n_{studies}	Mean Macro-F1
Orthopaedics & trauma	7	0.680
Endocrine & metabolic	10	0.638
Rheumatology	5	0.637
Wounds	5	0.628
Dentistry & oral health	15	0.597
Lungs & airways	14	0.589
Pregnancy & childbirth	9	0.583
Neonatal	17	0.578
Complementary & alternative medicine	19	0.555
Child health	72	0.550
Gastroenterology & hepatology	21	0.538
Heart & circulation	45	0.536
Neurology	11	0.519
Cancer	21	0.494
Effective practice & health systems	4	0.487
Kidney disease	31	0.484
Mental health	11	0.481
Gynaecology	11	0.480
Insurance medicine	21	0.477
Blood disorders	6	0.476
Genetic disorders	9	0.473
Skin disorders	3	0.448
Reproductive & sexual health	8	0.437
Infectious disease	8	0.425

model). The low ICC indicates that whether a model correctly judges one domain bears little information about its judgement on another domain of the same study, which is consistent with the domain-specific nature of each RoB criterion.

Using the Kish (1965) design-effect formula, the effective sample size is

$$n_{\text{eff}} = \frac{n_{\text{cells}}}{1 + (k - 1)\hat{\rho}} = \frac{1,589}{1 + 6 \times 0.044} \approx 1,257, \quad (\text{S2})$$

where $k = 7$ domains per study.

Table S6: Pairwise complementarity matrix. Cell (i, j) = % of 1,589 cells where row model i is incorrect and column model j is correct. Row model abbreviations: Haiku = Claude Haiku 4.5; GLM-5 = GLM-5; Qwen = Qwen 3.7-Plus; Flash = Gemini 2.5 Flash; DSV4 = DeepSeek v4-Flash; Grok = Grok 4.3; GPT = GPT-5.4; Pro = Gemini 2.5 Pro.

	Haiku	GLM-5	Qwen	Flash	DSV4	Grok	GPT	Pro
Haiku	—	12.5	10.6	9.9	12.3	12.5	12.6	12.4
GLM-5	13.9	—	7.4	12.3	9.6	8.9	9.5	7.6
Qwen	12.3	7.6	—	11.7	10.3	7.6	9.6	9.4
Flash	11.7	12.8	11.9	—	11.3	10.2	9.4	9.1
DSV4	14.5	10.4	10.8	11.7	—	9.6	9.5	9.9
Grok	14.9	9.9	8.4	10.8	9.8	—	9.0	8.4
GPT	17.4	12.9	12.7	12.4	12.1	11.4	—	10.1
Pro	17.3	11.1	12.6	12.2	12.6	11.0	10.3	—

1180 Numerical comparison

1181 Table S7 reports accuracy CIs for five models under three methods.

Table S7: Comparison of 95% confidence interval methods for overall accuracy. Width in percentage points (pp). Bootstrap: study-clustered, 5,000 iterations. Wilson n_{cells} : Equation (S1) with $n = 1,589$. Wilson n_{eff} : Equation (S1) with $n = 1,257$ (Kish formula). Wilson n_{studies} : Equation (S1) with $n = 227$.

Model	Accuracy	Wilson n_{cells} (width)	Wilson n_{eff} (width)	Bootstrap (width)
Claude Haiku 4.5	62.5%	[60.1, 64.8] (4.8)	[59.8, 65.1] (5.3)	[59.8, 65.2] (5.4)
Gemini 2.5 Flash	60.7%	[58.2, 63.0] (4.8)	[57.9, 63.4] (5.4)	[58.2, 63.1] (4.9)
DeepSeek v4-Flash	60.3%	[57.9, 62.7] (4.8)	[57.6, 63.0] (5.4)	[57.7, 62.9] (5.2)
GPT-5.4	57.7%	[55.3, 60.1] (4.9)	[54.9, 60.4] (5.4)	[55.1, 60.3] (5.2)
Gemini 2.5 Pro	57.6%	[55.1, 60.0] (4.9)	[54.8, 60.3] (5.4)	[55.1, 60.0] (4.9)
Wilson $n_{\text{studies}} = 227$		widths ≈ 12.5 – 12.8 pp (not shown; see text)		

1182 Wilson with n_{cells} underestimates uncertainty by $\sim 10\%$ relative to bootstrap
1183 (widths 4.8 pp vs. 5.1 pp). Wilson with n_{eff} virtually matches the bootstrap (widths
1184 5.3 pp vs. 5.1 pp); the two methods agree to within 0.1 pp for every model. Wilson
1185 with n_{studies} overestimates uncertainty by $\sim 2.5\times$ because it discards the domain-
1186 level information and treats each study as a single binary observation.

1187 Why bootstrap was adopted for all metrics

1188 The numerical equivalence between Wilson n_{eff} and bootstrap for accuracy con-
1189 firms that the low ICC ($\hat{\rho} = 0.044$) makes the independence assumption nearly
1190 tenable. Nevertheless, bootstrap was used uniformly for three reasons. First,
1191 Macro-F1 and Gwet's AC1 are composite statistics with no closed-form variance ex-
1192 pression; bootstrap applies without modification. Second, the paired-bootstrap
1193 procedure used for inference-mode Δ values (Table 6) jointly resamples both
1194 conditions, correctly propagating correlation between them; Wilson intervals
1195 cannot be directly adapted for paired differences of non-proportion statistics.
1196 Third, a single method for all reported CIs ensures internal consistency and
1197 simplifies interpretation.

S8 Inference Cost

Table S8 reports total billed inference costs for the eight models across the 227-study benchmark, retrieved from the OpenRouter account history. All costs correspond to per-domain inference mode (seven sequential calls per study). The total cost of running the full eight-model benchmark was approximately \$244 USD, providing an order-of-magnitude reference for the practical feasibility of this type of evaluation.

The pipeline permitted up to two retries per domain call in the event of API errors, timeouts, or invalid responses, and some studies required full re-queuing; the number of retry attempts was not recorded per model. Retry attempts were sufficient to ensure that all 1,589 study-domain cells received a valid prediction for every model: no cells were excluded from analysis due to unresolvable output failures. Billed costs therefore reflect total API consumption including failed calls, so absolute values should not be used to derive a cost-per-assessment figure for any individual model. The *ordering* of billed costs across models is nonetheless informative: because input tokens were identical across models (same corpus and prompt) and output token counts are of similar magnitude, differences in list price per token dominate over retry-rate variation. The observed cost ordering—DeepSeek v4-Flash cheapest, GPT-5.4 most expensive—is fully consistent with published list prices on OpenRouter (Table SS8), confirming that price per token is the primary driver of cost differences in this benchmark.

For Gemini 2.5 Flash and DeepSeek v4-Flash, which were also evaluated in single-call mode (Section 3.5), billed single-call costs were lower than per-domain costs (\$2.05 vs. \$9.42 and \$0.54 vs. \$3.50 respectively). The direction of this difference is expected—single-call issues one API request per study versus seven—but the magnitude should not be interpreted as a precise cost ratio because retry counts were not recorded for either inference mode.

The costs reported here reflect API-based access via OpenRouter and apply specifically to proprietary or API-only model endpoints. Several models evalu-

Table S8: Total billed inference costs for the eight models (per-domain mode, 227 studies), retrieved from OpenRouter account history, 7 July 2026. Costs include retries due to API errors or invalid responses; the number of retries was not recorded per model and cross-model comparisons should not be drawn from these figures. Total across all eight models: \$244.32.

Model	Total billed cost (USD)
DeepSeek v4-Flash	3.50
Gemini 2.5 Flash	9.42
Qwen 3.7-Plus	11.40
GLM-5	30.70
Claude Haiku 4.5	32.40
Grok 4.3	33.50
Gemini 2.5 Pro	54.10
GPT-5.4	69.30
Total	244.32

1227

1228

1229

1230

1231

1232

1233

1234

1235

1236

1237

1238

1239

1240

1241

1242

1243

1244

ated in this benchmark belong to families that release open weights publicly. DeepSeek v4-Flash weights are publicly available under an MIT licence on Hugging Face (`deepseek-ai/DeepSeek-V4-Flash`), and Nemotron-3-Ultra-550B weights are similarly available under the OpenMDW-1.1 licence (`nvidia/NVIDIA-Nemotron-3-Ultra-550B-A55B`). both models could therefore be self-hosted at no per-inference monetary cost by researchers with access to sufficient GPU resources. The situation for GLM-5 is more nuanced: the specific checkpoint served via the `z-ai/glm-5` OpenRouter endpoint at the time of evaluation was API-only, but Zhipu AI subsequently published weights for GLM-5.1 to the same Hugging Face repository (`zai-org/GLM-5`); the exact correspondence between the API-served and the released checkpoint is not confirmed by the provider. Qwen 3.7-Plus was API-only at the time of evaluation (May–July 2026); Alibaba has historically released open weights for smaller Qwen variants but had not published weights for the Plus tier as of the evaluation window. The remaining five models—Gemini 2.5 Pro, Gemini 2.5 Flash, Claude Haiku 4.5, GPT-5.4, and Grok 4.3—are proprietary and API-only with no publicly released weights, and their inference costs are therefore irreducible below the API pricing reported in Table S8. The choice of API-based access in the present study was motivated by reproducibility under identical conditions

across all eight models and by the absence of hardware capable of serving
550B-parameter models locally.

S9 Expanded Formulas for Cohen's κ and Gwet's AC1

Let $\mathbf{N} = (N_{ij})$ be the $K \times K$ confusion matrix ($K = 3$), with rows indexed by
predicted class i and columns indexed by ground-truth class j , $N = \sum_{i,j} N_{ij}$ the
total number of evaluated cells, and

$$\hat{p}_{k,\text{model}} = \frac{1}{N} \sum_{j=1}^K N_{kj}, \quad \hat{p}_{k,\text{gt}} = \frac{1}{N} \sum_{i=1}^K N_{ik}$$

the marginal proportions of class k in the model predictions and ground truth
(row and column sums of $\mathbf{N}$ divided by N), as defined in the Methods.

Cohen's κ . The chance-expected agreement is

$$p_e^\kappa = \sum_{k=1}^K \hat{p}_{k,\text{model}} \cdot \hat{p}_{k,\text{gt}},$$

where each term $\hat{p}_{k,\text{model}} \cdot \hat{p}_{k,\text{gt}}$ is the probability that both the model and ground
truth independently assign class k to a randomly selected cell. Substituting
into $\kappa = (p_o - p_e^\kappa)/(1 - p_e^\kappa)$ with $p_o = \sum_k N_{kk}/N$ gives the equivalent form in raw
counts:

$$\kappa = \frac{N \sum_{k=1}^K N_{kk} - \sum_{k=1}^K \left(\sum_{j=1}^K N_{kj} \right) \left(\sum_{i=1}^K N_{ik} \right)}{N^2 - \sum_{k=1}^K \left(\sum_{j=1}^K N_{kj} \right) \left(\sum_{i=1}^K N_{ik} \right)}.$$

Gwet's AC1. The mean marginal proportion for class k is

$$\pi_k = \frac{\hat{p}_{k,\text{model}} + \hat{p}_{k,\text{gt}}}{2},$$

1259 and the chance-agreement term is

$$p_e = \frac{1}{K-1} \sum_{k=1}^K \pi_k (1 - \pi_k).$$

1260 $AC1 = (p_o - p_e)/(1 - p_e)$ follows with p_o as above.

1261 The key difference between the two coefficients lies in their chance terms: p_e^k
1262 depends on the *product* of the two marginals, making it sensitive to imbalance
1263 in either rater’s distribution; p_e for AC1 depends on the *mean* marginal π_k , which
1264 is more stable when one rater (here, the ground truth) has highly skewed class
1265 proportions.

1266 **S10 Prompt Used for Figure Description (VLM Pipeline)**

1267 During Stage 2 of the document processing pipeline (Section 2.2), each figure
1268 placeholder was submitted to Gemini 2.5 Flash (google/gemini-2.5-flash) with
1269 the following fixed prompt. The system message was:

1270 You are converting a scientific figure into faithful Markdown text.
1271 Describe only what is visibly present in the figure.
1272 Do not infer unstated conclusions, do not mention OCR or the model,
1273 and do not use JSON.
1274 Return Markdown only.

1275 The user message was:

1276 Rewrite this scientific figure as Markdown so the image can be
1277 removed from the extracted article.

1278

1279 Requirements:

- 1280 - Preserve scientific meaning.
- 1281 - Identify figure label and caption when visible or provided in context.
- 1282 - Describe chart/figure type.
- 1283 - Describe axes, units, legend, groups, markers, and error bars when visible.

- 1284 - Summarize visible trends and approximate plotted relationships without
- 1285 inventing exact values.
- 1286 - Mention unreadable or ambiguous areas explicitly.
- 1287 - Keep the output concise but detailed enough to replace the image in
- 1288 the article.

1289

1290 Use this Markdown template:

1291 **### Figure Description**

- 1292 - source_image: '...'
- 1293 - figure_label: ...
- 1294 - caption: ...
- 1295 - figure_type: ...
- 1296 - axes_and_units: ...
- 1297 - legend_and_groups: ...
- 1298 - visual_summary: ...
- 1299 - plotted_patterns: ...
- 1300 - unreadable_or_uncertain: ...

1301

1302 Nearby article context:

1303 {context}

1304

1305 Image path reference: {image_path}

1306 where {context} is the paragraph immediately surrounding the figure in the
1307 extracted Markdown, and {image_path} is the relative path to the page image
1308 extracted by MinerU. The prompt was used without modification for all figures
1309 across all 227 studies.

1310 **S11 Stratification by Publication Era**

1311 To investigate whether study characteristics associated with publication period
1312 affect model performance, the 224 studies with a recorded publication year were

grouped into four eras: Pre-2000s (1960–1999; $n = 32$), 2000s (2000–2009; $n = 44$), 2010s (2010–2019; $n = 105$), and 2020s (2020 onward; $n = 43$). Figure S3 shows the distribution of studies by publication decade. Figures S4 and S5 show overall accuracy per era aggregated across all eight models and per model, respectively. Figure S6 shows accuracy per RoB domain across eras. All confidence intervals are 95% study-clustered bootstrap (5,000 iterations).

Point estimates suggest a modest trend of increasing accuracy from Pre-2000s (56.0%) to the 2020s (63.9%); however, confidence intervals overlap across all pairwise comparisons, including the most extreme (Pre-2000s vs. 2020s: [50.2, 61.7] vs. [58.7, 69.3]), so a monotonic trend cannot be formally established. The wide interval for the Pre-2000s era reflects the small stratum size ($n = 32$) rather than greater variability in model performance.

Domain-level stratification reveals heterogeneous patterns. D1 (random sequence generation) shows the clearest era effect, rising from 71% in the Pre-2000s to 88% in the 2020s, consistent with the progressive adoption of CONSORT-mandated reporting of randomisation methods. D3 (blinding of participants and personnel) shows the opposite trend, declining from 64% in the Pre-2000s to 53% in the 2020s, which may reflect the increasing prevalence of non-pharmacological and device trials in recent decades, where blinding is structurally more complex to evaluate. D7 (other bias) remains the lowest-performing domain across all eras (38–45%), with no discernible temporal trend, consistent with the domain reflecting intrinsic criterion vagueness rather than a reporting-period effect.




Figure S3: Distribution of 227 RCTs by publication decade. Numbers above each bar indicate study count per decade.




Figure S4: Overall accuracy by publication era, aggregated across all eight models (95% CI, study-clustered bootstrap, 5,000 iterations). Point estimates are annotated above the upper bound of each confidence band. The Pre-2000s stratum contains only 32 studies, yielding a substantially wider interval than the remaining three eras.




Figure S5: Overall accuracy by publication era for each of the eight models individually (95% CI, study-clustered bootstrap, 5,000 iterations). Shaded bands show the confidence interval for each model. The two-tier structure observed in global performance (Gemini 2.5 Pro and GPT-5.4 below the remaining six models) is consistent across all four eras.




Figure S6: Accuracy per RoB domain across publication eras, aggregated across all eight models (95% CI, study-clustered bootstrap, 5,000 iterations). D1 (sequence generation) shows a pronounced positive trend; D3 (blinding of participants and personnel) shows a negative trend; D7 (other bias) remains the lowest domain in all eras with no clear temporal pattern.